

Lecture 3: Data Wrangling I

Data Science for Business Analytics

Thibault Vatter <thibault.vatter@unil.ch>

Department of Statistics, Columbia University and HEC Lausanne, UNIL

12.03.2018

1 Overview

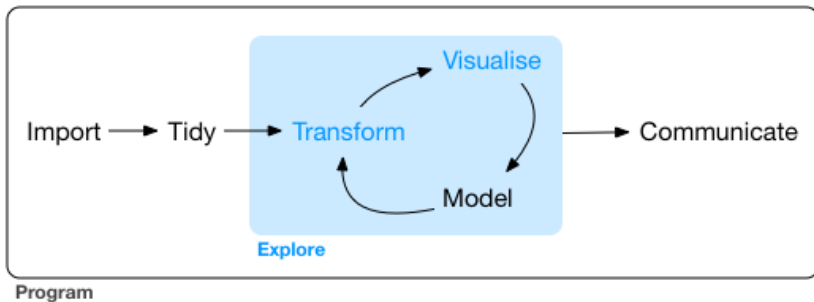
2 Tibbles

3 Data manipulation

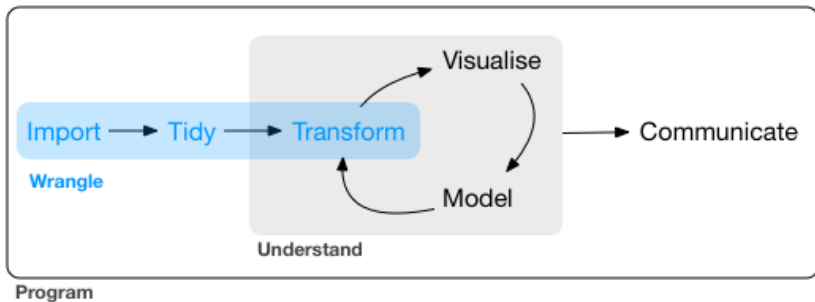
4 The %>% operator

5 More on data manipulation

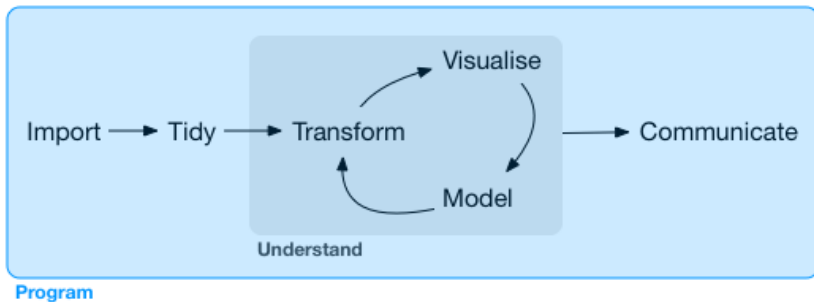
6 Tidy data



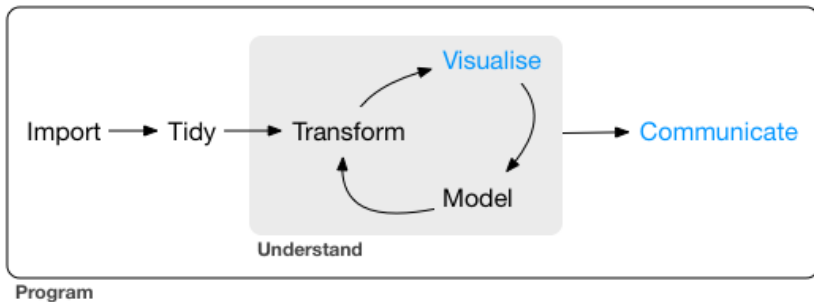
source: R for Data Science (like most figures in what follows)



This morning



This afternoon



1 Overview

2 Tibbles

3 Data manipulation

4 The %>% operator

5 More on data manipulation

6 Tidy data

- Alternative R's traditional `data.frame`.
- Tweak some older behaviours to make life easier.
- Part of the core tidyverse.
- Unifying feature of the tidyverse.
- Most functions from the tidyverse produce tibbles.
- To learn more, see `vignette("tibble")`.

Coerce a data frame to a tibble

```
as_tibble(iris)
```

```
## # A tibble: 150 x 5
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
##   <dbl>         <dbl>         <dbl>         <dbl> <fct>
## 1         5.10         3.50         1.40         0.200 setosa
## 2         4.90         3.00         1.40         0.200 setosa
## 3         4.70         3.20         1.30         0.200 setosa
## 4         4.60         3.10         1.50         0.200 setosa
## 5         5.00         3.60         1.40         0.200 setosa
## 6         5.40         3.90         1.70         0.400 setosa
## 7         4.60         3.40         1.40         0.300 setosa
## 8         5.00         3.40         1.50         0.200 setosa
## 9         4.40         2.90         1.40         0.200 setosa
## 10        4.90         3.10         1.50         0.100 setosa
## # ... with 140 more rows
```

Create from individual vectors

```
tibble(x = 1:5,  
y = 1,  
z = x ^ 2 + y)
```

```
## # A tibble: 5 x 3  
##       x     y     z  
##   <int> <dbl> <dbl>  
## 1     1  1.00  2.00  
## 2     2  1.00  5.00  
## 3     3  1.00 10.0  
## 4     4  1.00 17.0  
## 5     5  1.00 26.0
```

```
tribble(  
  ~colA, ~colB,  
  "a",   1,  
  "b",   2,  
  "c",   3  
)
```

```
## # A tibble: 3 x 2  
##   colA    colB  
##   <chr> <dbl>  
## 1 a      1.00  
## 2 b      2.00  
## 3 c      3.00
```

```
(df <- tibble(a = lubridate::today() + runif(4e1) * 30,  
b = 1:4e1,  
c = runif(4e1),  
d = sample(letters, 4e1, replace = TRUE)))
```

```
## # A tibble: 40 x 4  
##       a           b         c d  
##   <date>    <int>  <dbl> <chr>  
## 1 2018-04-01      1 0.411  y  
## 2 2018-03-13      2 0.821  l  
## 3 2018-03-17      3 0.647  s  
## 4 2018-03-23      4 0.783  k  
## 5 2018-04-02      5 0.553  i  
## 6 2018-03-12      6 0.530  t  
## 7 2018-04-01      7 0.789  f  
## 8 2018-04-03      8 0.0233 s  
## 9 2018-03-25      9 0.477  d  
## 10 2018-03-24     10 0.732  g  
## # ... with 30 more rows
```

```
print(df, n = 2, width = 30)
```

```
## # A tibble: 40 x 4
##   a           b       c
##   <date>     <int> <dbl>
## 1 2018-04-01     1 0.411
## 2 2018-03-13     2 0.821
## # ... with 38 more rows, and
## #   1 more variable: d <chr>
```

- `options(tibble.print_max = n, tibble.print_min = m)`: if more than `m` rows, print only `n` rows.
- `options(dplyr.print_min = Inf)` to always show all rows.
- `options(tibble.width = Inf)` to always print all columns.
- `package?tibble`

```
# Extract by name I  
df$b
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
## [24] 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

```
# Extract by name II  
df[["b"]]
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
## [24] 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

```
# Extract by position  
df[[2]]
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
## [24] 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

Compared to a `data.frame`:

- no partial matching
- warning if the column does not exist

1 Overview

2 Tibbles

3 Data manipulation

4 The %>% operator

5 More on data manipulation

6 Tidy data

When working with data you must:

- Figure out what you want to do.
- Describe those tasks in the form of a computer program.
- Execute the program.

dplyr makes these steps fast and easy:

- By constraining your options, it helps you think about your data manipulation challenges.
- It provides simple **“verbs”**, functions that correspond to the most common data **manipulation tasks**, to help you translate your thoughts into code.
- It uses efficient backends, so you spend less time waiting for the computer.

5 verbs to solve common data manipulation challenges:

- **`filter()` to select observations** based on their values.
- **`arrange()` to reorder the observations.**
- **`select()` to select variables** based on their names.
- **`mutate()` to add variables** as functions of existing variables.
- **`summarize()` to collapse many values** down to a single summary.

Two important features:

- All verbs operate groupwise with `group_by()`.
- All verbs work similarly:
 1. First argument is a data frame.
 2. Other arguments describe what to do with it using variable names.
 3. Result is a new data frame.

All 336,776 flights that departed from NYC in 2013 (US BTS):

```
nycflights13::flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
##  1  2013     1     1     517           515         2.00     830
##  2  2013     1     1     533           529         4.00     850
##  3  2013     1     1     542           540         2.00     923
##  4  2013     1     1     544           545        -1.00    1004
##  5  2013     1     1     554           600        -6.00     812
##  6  2013     1     1     554           558        -4.00     740
##  7  2013     1     1     555           600        -5.00     913
##  8  2013     1     1     557           600        -3.00     709
##  9  2013     1     1     557           600        -3.00     838
## 10  2013     1     1     558           600        -2.00     753
## # ... with 336,766 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

Filter rows with filter()

```
filter(flights, month == 1, day == 1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515        2.00     830
## 2  2013     1     1     533           529        4.00     850
## 3  2013     1     1     542           540        2.00     923
## 4  2013     1     1     544           545       -1.00    1004
## 5  2013     1     1     554           600       -6.00     812
## 6  2013     1     1     554           558       -4.00     740
## 7  2013     1     1     555           600       -5.00     913
## 8  2013     1     1     557           600       -3.00     709
## 9  2013     1     1     557           600       -3.00     838
## 10 2013     1     1     558           600       -2.00     753
## # ... with 832 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

- The standard suite: $>$, $>=$, $<$, $<=$, $!=$, and $==$.
- Most common mistake:

```
filter(flights, month = 1)
```


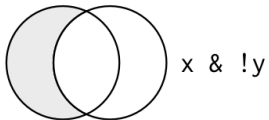
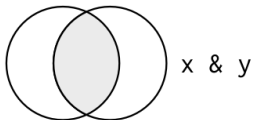
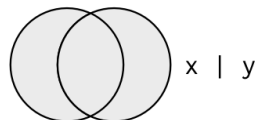

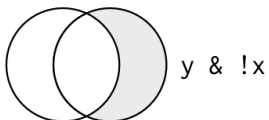
- What happens in the following?

```
sqrt(2) ^ 2 == 2  
1/49 * 49 == 1  
near(sqrt(2) ^ 2, 2)  
near(1 / 49 * 49, 1)
```

```
## [1] FALSE  
## [1] FALSE  
## [1] TRUE  
## [1] TRUE
```

Multiple arguments to `filter()` are combined with:

- `&` for “and”
- `|` for “or”
- `!` for “not”



What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

Literally “finds all flights that departed in November or December”, but you can’t write `filter(flights, month == 11 | 12)`.

What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

Literally “finds all flights that departed in November or December”, but you can’t write `filter(flights, month == 11 | 12)`.
Solution:

```
filter(flights, month %in% c(11, 12))
```


- $!(x \ \& \ y)$ is the same as $!x \mid !y$
- $!(x \mid y)$ is the same as $!x \ \& \ !y$

```
filter(flights, !(arr_delay > 120 | dep_delay > 120))  
filter(flights, arr_delay <= 120, dep_delay <= 120)
```

NAs (“not availables”) are “contagious”:

```
NA > 5  
10 == NA  
NA + 10  
NA / 2  
NA == NA
```

```
## [1] NA  
## [1] NA  
## [1] NA  
## [1] NA  
## [1] NA
```

To determine if a value is missing‘:

```
is.na(NA)
```

```
## [1] TRUE
```

Missing values and filter()

```
df <- tibble(x = c(1, NA, 3))
```

```
filter(df, x > 1)
```

```
## # A tibble: 1 x 1
##       x
##   <dbl>
## 1  3.00
```

```
filter(df, is.na(x) | x > 1)
```

```
## # A tibble: 2 x 1
##       x
##   <dbl>
## 1  NA
## 2  3.00
```

Arrange rows with arrange()

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517             515         2.00     830
## 2  2013     1     1     533             529         4.00     850
## 3  2013     1     1     542             540         2.00     923
## 4  2013     1     1     544             545        -1.00    1004
## 5  2013     1     1     554             600        -6.00     812
## 6  2013     1     1     554             558        -4.00     740
## 7  2013     1     1     555             600        -5.00     913
## 8  2013     1     1     557             600        -3.00     709
## 9  2013     1     1     557             600        -3.00     838
## 10 2013     1     1     558             600        -2.00     753
## # ... with 336,766 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

arrange() and desc()

```
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9     641             900         1301    1242
## 2  2013     6    15    1432            1935         1137    1607
## 3  2013     1    10    1121            1635         1126    1239
## 4  2013     9    20    1139            1845         1014    1457
## 5  2013     7    22     845            1600         1005    1044
## 6  2013     4    10    1100            1900          960    1342
## 7  2013     3    17    2321             810          911     135
## 8  2013     7    22    2257             759          898     121
## 9  2013    12     5     756            1700          896   1058
## 10 2013     5     3    1133            2055          878   1250
## # ... with 336,766 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

arrange() and missing values

```
df <- tibble(x = c(5, NA, 2))  
arrange(df, x)
```

```
## # A tibble: 3 x 1  
##       x  
##   <dbl>  
## 1  2.00  
## 2  5.00  
## 3 NA
```

```
arrange(df, desc(x))
```

```
## # A tibble: 3 x 1  
##       x  
##   <dbl>  
## 1  5.00  
## 2  2.00  
## 3 NA
```

Select columns with `select()`

```
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

All columns between year and day

```
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```


All columns except from year to day

```
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int>         <int>      <dbl>   <int>         <int>
## 1      517           515        2.00     830           819
## 2      533           529        4.00     850           830
## 3      542           540        2.00     923           850
## 4      544           545       -1.00    1004          1022
## 5      554           600       -6.00     812           837
## 6      554           558       -4.00     740           728
## 7      555           600       -5.00     913           854
## 8      557           600       -3.00     709           723
## 9      557           600       -3.00     838           846
## 10     558           600       -2.00     753           745
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>,
## #   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

select() and everything()

```
select(flights, time_hour, air_time, everything())
```

```
## # A tibble: 336,776 x 19
##   time_hour          air_time year month   day dep_time
##   <dtm>            <dbl> <int> <int> <int>   <int>
## 1 2013-01-01 05:00:00    227   2013     1     1     517
## 2 2013-01-01 05:00:00    227   2013     1     1     533
## 3 2013-01-01 05:00:00    160   2013     1     1     542
## 4 2013-01-01 05:00:00    183   2013     1     1     544
## 5 2013-01-01 06:00:00    116   2013     1     1     554
## 6 2013-01-01 05:00:00    150   2013     1     1     554
## 7 2013-01-01 06:00:00    158   2013     1     1     555
## 8 2013-01-01 06:00:00     53.0 2013     1     1     557
## 9 2013-01-01 06:00:00    140   2013     1     1     557
## 10 2013-01-01 06:00:00    138   2013     1     1     558
## # ... with 336,766 more rows, and 13 more variables:
## #   sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   distance <dbl>, hour <dbl>, minute <dbl>
```

- Helper functions you can use within `select()`:
 - ▶ `starts_with("abc")`: matches names that begin with "abc".
 - ▶ `ends_with("xyz")`: matches names that end with "xyz".
 - ▶ `contains("ijk")`: matches names that contain "ijk".
 - ▶ `matches("(.)\\1")`: selects variables that match a regular expression (this one matches any variables that contain repeated characters).
 - ▶ `num_range("x", 1:3)` matches `x1`, `x2` and `x3`.
- `select()` can be used to rename variables, but it drops all of the variables not explicitly mentioned. Instead, use `rename()`
- See `?select` for more details.

Create a narrower dataset

```
(flights_sml <- select(flights,  
  year:day,  
  ends_with("delay"),  
  distance,  
  air_time))
```

```
## # A tibble: 336,776 x 7
```

```
##   year month   day dep_delay arr_delay distance air_time  
##   <int> <int> <int>    <dbl>    <dbl>    <dbl>    <dbl>  
## 1  2013     1     1      2.00     11.0     1400     227  
## 2  2013     1     1      4.00     20.0     1416     227  
## 3  2013     1     1      2.00     33.0     1089     160  
## 4  2013     1     1     -1.00    -18.0     1576     183  
## 5  2013     1     1     -6.00    -25.0      762     116  
## 6  2013     1     1     -4.00     12.0      719     150  
## 7  2013     1     1     -5.00     19.0     1065     158  
## 8  2013     1     1     -3.00    -14.0      229     53.0  
## 9  2013     1     1     -3.00     - 8.00      944     140  
## 10 2013     1     1     -2.00      8.00      733     138  
## # ... with 336,766 more rows
```

Add new variables with mutate()

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  speed = distance / air_time * 60)
```

```
## # A tibble: 336,776 x 9  
##   year month   day dep_delay arr_delay distance air_time   gain  
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>  
## 1  2013     1     1     2.00    11.0    1400    227     9.00  
## 2  2013     1     1     4.00    20.0    1416    227    16.0  
## 3  2013     1     1     2.00    33.0    1089    160    31.0  
## 4  2013     1     1    -1.00   -18.0    1576    183   -17.0  
## 5  2013     1     1    -6.00   -25.0     762    116  -19.0  
## 6  2013     1     1    -4.00    12.0     719    150    16.0  
## 7  2013     1     1    -5.00    19.0    1065    158    24.0  
## 8  2013     1     1    -3.00   -14.0     229     53.0 -11.0  
## 9  2013     1     1    -3.00    - 8.00     944    140    - 5.00  
## 10 2013     1     1    -2.00     8.00     733    138    10.0  
## # ... with 336,766 more rows, and 1 more variable: speed <dbl>
```

Refer to columns just created

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours)
```

```
## # A tibble: 336,776 x 10  
##   year month   day dep_delay arr_delay distance air_time  gain  
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>   <dbl> <dbl>  
## 1  2013     1     1       2.00      11.0      1400    227    9.00  
## 2  2013     1     1       4.00      20.0      1416    227   16.0  
## 3  2013     1     1       2.00      33.0      1089    160   31.0  
## 4  2013     1     1      -1.00     -18.0      1576    183  -17.0  
## 5  2013     1     1      -6.00     -25.0       762    116 -19.0  
## 6  2013     1     1      -4.00      12.0       719    150   16.0  
## 7  2013     1     1      -5.00      19.0      1065    158   24.0  
## 8  2013     1     1      -3.00     -14.0       229    53.0 -11.0  
## 9  2013     1     1      -3.00     - 8.00       944    140   - 5.00  
## 10 2013     1     1      -2.00      8.00       733    138   10.0  
## # ... with 336,766 more rows, and 2 more variables: hours <dbl>,  
## #   gain_per_hour <dbl>
```

```
transmute(flights,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours)
```

```
## # A tibble: 336,776 x 3  
##       gain hours gain_per_hour  
##   <dbl> <dbl>      <dbl>  
## 1    9.00  3.78         2.38  
## 2   16.0   3.78         4.23  
## 3   31.0   2.67        11.6  
## 4  -17.0   3.05        - 5.57  
## 5  -19.0   1.93        - 9.83  
## 6   16.0   2.50         6.40  
## 7   24.0   2.63         9.11  
## 8  -11.0  0.883        -12.5  
## 9   - 5.00  2.33        - 2.14  
## 10  10.0   2.30         4.35  
## # ... with 336,766 more rows
```

Any vectorized function would work, but frequently useful are:

- Arithmetic operators: $+$, $-$, $*$, $/$, $^$.
 - ▶ Vectorized with “recycling rules” (e.g., `air_time / 60`).
 - ▶ Useful in conjunction with aggregate functions (e.g., `x / sum(x)` or `y - mean(y)`).
- Modular arithmetic: `%/%` (integer division) and `%%` (remainder), where `x == y * (x %/% y) + (x %% y)`.
 - ▶ Allows you to break integers up into pieces (e.g., `hour = dep_time %/% 100` and `minute = dep_time %% 100`)
- Logs: `log()`, `log2()`, `log10()`.
 - ▶ Useful for data ranging across multiple orders of magnitude.
 - ▶ Convert multiplicative relationships to additive.

■ Offsets: `lead()` and `lag()`:

- ▶ Refer to lead-/lagging values (e.g. to get running differences $x - \text{lag}(x)$ or find when values change $x \neq \text{lag}(x)$).
- ▶ Useful in conjunction with `group_by()`.

```
x <- 1:10  
lag(x)  
lead(x)
```

```
## [1] NA  1  2  3  4  5  6  7  8  9  
## [1]  2  3  4  5  6  7  8  9 10 NA
```

■ Cumulative aggregates: `cumsum()`, `cumprod()`, `cummin()`, `cummax()`, `cummean()` (RcppRoll package for rolling aggregates).

```
cumsum(x)  
cummean(x)
```

```
## [1]  1  3  6 10 15 21 28 36 45 55  
## [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5
```

- Logical comparisons, <, <=, >, >=, !=
- Ranking functions: `min_rank()`, `row_number()`, `dense_rank()`, `percent_rank()`, `cume_dist()`, `ntile()`

```
y <- c(1, 2, 2, NA, 3, 4)
min_rank(y)
min_rank(desc(y))
row_number(y)
dense_rank(y)
percent_rank(y)
cume_dist(y)
```

```
## [1] 1 2 2 NA 4 5
## [1] 5 3 3 NA 2 1
## [1] 1 2 3 NA 4 5
## [1] 1 2 2 NA 3 4
## [1] 0.00 0.25 0.25 NA 0.75 1.00
## [1] 0.2 0.6 0.6 NA 0.8 1.0
```

Collapse values with `summarize()`

```
summarize(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 1
##   delay
##   <dbl>
## 1  12.6
```

summarize() paired with group_by()

```
by_day <- group_by(flights, year, month, day)
summarize(by_day, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [?]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
## 9  2013     1     9   2.28
## 10 2013     1    10   2.84
## # ... with 355 more rows
```

- 1 Overview
- 2 Tibbles
- 3 Data manipulation
- 4 The %>% operator**
- 5 More on data manipulation
- 6 Tidy data

What is this code doing?

```
a1 <- group_by(flights, year, month, day)
a2 <- select(a1, arr_delay, dep_delay)
a3 <- summarize(a2,
                arr = mean(arr_delay, na.rm = TRUE),
                dep = mean(dep_delay, na.rm = TRUE))
filter(a3, arr > 30 | dep > 30)
```

```
## # A tibble: 49 x 5
## # Groups:   year, month [11]
##   year month   day   arr   dep
##   <int> <int> <int> <dbl> <dbl>
## 1  2013     1    16  34.2  24.6
## 2  2013     1    31  32.6  28.7
## 3  2013     2    11  36.3  39.1
## 4  2013     2    27  31.3  37.8
## 5  2013     3     8  85.9  83.5
## 6  2013     3    18  41.3  30.1
## 7  2013     4    10  38.4  33.0
## 8  2013     4    12  36.0  34.8
## 9  2013     4    18  36.0  34.9
## 10 2013     4    19  47.9  46.1
## # ... with 39 more rows
```

Same code (no unnecessary objects)

```
filter(summarize(select(group_by(flights, year, month, day),
  arr_delay, dep_delay),
  arr = mean(arr_delay, na.rm = TRUE),
  dep = mean(dep_delay, na.rm = TRUE)),
  arr > 30 | dep > 30)
```

```
## # A tibble: 49 x 5
## # Groups:   year, month [11]
##   year month   day   arr   dep
##   <int> <int> <int> <dbl> <dbl>
## 1  2013     1    16  34.2  24.6
## 2  2013     1    31  32.6  28.7
## 3  2013     2    11  36.3  39.1
## 4  2013     2    27  31.3  37.8
## 5  2013     3     8  85.9  83.5
## 6  2013     3    18  41.3  30.1
## 7  2013     4    10  38.4  33.0
## 8  2013     4    12  36.0  34.8
## 9  2013     4    18  36.0  34.9
## 10 2013     4    19  47.9  46.1
## # ... with 39 more rows
```

... Or use %>%

```
flights %>%  
  group_by(year, month, day) %>%  
  select(arr_delay, dep_delay) %>%  
  summarize(arr = mean(arr_delay, na.rm = TRUE),  
            dep = mean(dep_delay, na.rm = TRUE)) %>%  
  filter(arr > 30 | dep > 30)
```

```
## # A tibble: 49 x 5  
## # Groups:   year, month [11]  
##   year month   day   arr   dep  
##   <int> <int> <int> <dbl> <dbl>  
## 1  2013     1    16  34.2  24.6  
## 2  2013     1    31  32.6  28.7  
## 3  2013     2    11  36.3  39.1  
## 4  2013     2    27  31.3  37.8  
## 5  2013     3     8  85.9  83.5  
## 6  2013     3    18  41.3  30.1  
## 7  2013     4    10  38.4  33.0  
## 8  2013     4    12  36.0  34.8  
## 9  2013     4    18  36.0  34.9  
## 10 2013     4    19  47.9  46.1  
## # ... with 39 more rows
```


Makes your code more readable by:

- structuring sequences of data operations left-to-right,
- minimizing the need for local variables and function definitions,
- making it easy to add steps anywhere in the sequence of operations.

- `x %>% f` is equivalent to `f(x)`
- `x %>% f(y)` is equivalent to `f(x, y)`
- `x %>% f(y) %>% g(z)` is equivalent to `g(f(x, y), z)`

```
x <- 1:10  
y <- x + 1  
z <- y + 1  
f <- function(x, y) x + y
```

```
x %>% sum
```

```
## [1] 55
```

```
x %>% f(y)
```

```
## [1] 3 5 7 9 11 13 15 17 19 21
```

```
x %>% f(y) %>% f(z)
```

```
## [1] 6 9 12 15 18 21 24 27 30 33
```

The argument (“dot”) placeholder

- `x %>% f(y, .)` is equivalent to `f(y, x)`
- `x %>% f(y, z = .)` is equivalent to `f(y, z = x)`

```
x <- 1:10
y <- 2 * x
f <- function(z, y) y / z

x %>% f(y, .)
```

```
## [1] 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5
```

```
x %>% f(y, z = .)
```

```
## [1] 2 2 2 2 2 2 2 2 2 2
```

Subsetting tibbles revisited

```
df <- tibble(a = lubridate::today() + runif(4e1) * 30,  
             b = 1:4e1,  
             c = runif(4e1),  
             d = sample(letters, 4e1, replace = TRUE))
```

```
# Extract by name  
df %>% .$b
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
## [24] 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

```
# Extract by position  
df %>% .[["b"]]
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23  
## [24] 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

- 1 Overview
- 2 Tibbles
- 3 Data manipulation
- 4 The %>% operator
- 5 More on data manipulation**
- 6 Tidy data

What is happening here?

```
flights %>%  
  group_by(year, month, day) %>%  
  summarize(mean = mean(dep_delay))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1    NA  
## 2  2013     1     2    NA  
## 3  2013     1     3    NA  
## 4  2013     1     4    NA  
## 5  2013     1     5    NA  
## 6  2013     1     6    NA  
## 7  2013     1     7    NA  
## 8  2013     1     8    NA  
## 9  2013     1     9    NA  
## 10 2013     1    10    NA  
## # ... with 355 more rows
```

Use `na.rm = TRUE`

```
flights %>%  
  group_by(year, month, day) %>%  
  summarize(mean = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day   mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1  11.5  
## 2  2013     1     2  13.9  
## 3  2013     1     3  11.0  
## 4  2013     1     4   8.95  
## 5  2013     1     5   5.73  
## 6  2013     1     6   7.15  
## 7  2013     1     7   5.42  
## 8  2013     1     8   2.55  
## 9  2013     1     9   2.28  
## 10 2013     1    10   2.84  
## # ... with 355 more rows
```

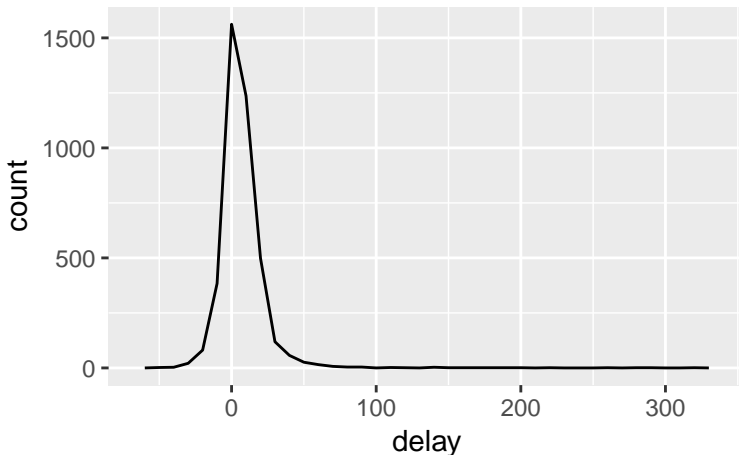
Or pre-filter the dataset

```
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))  
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(mean = mean(dep_delay))
```


```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1  11.4  
## 2  2013     1     2  13.7  
## 3  2013     1     3  10.9  
## 4  2013     1     4   8.97  
## 5  2013     1     5   5.73  
## 6  2013     1     6   7.15  
## 7  2013     1     7   5.42  
## 8  2013     1     8   2.56  
## 9  2013     1     9   2.30  
## 10 2013     1    10   2.84  
## # ... with 355 more rows
```


What do you see?

```
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarize(delay = mean(arr_delay))
```



```
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarize(delay = mean(arr_delay, na.rm = TRUE), n = n())
```



- Measures of location: `mean()`, `median()`.
- Measures of spread: `sd()`, `IQR()`, `mad()`.

- Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`.

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(first = min(dep_time), last = max(dep_time))
```

```
## # A tibble: 365 x 5  
## # Groups:   year, month [?]  
##   year month   day first last  
##   <int> <int> <int> <dbl> <dbl>  
## 1  2013     1     1  517  2356  
## 2  2013     1     2  42.0  2354  
## 3  2013     1     3  32.0  2349  
## 4  2013     1     4  25.0  2358  
## 5  2013     1     5  14.0  2357  
## 6  2013     1     6  16.0  2355  
## 7  2013     1     7  49.0  2359  
## 8  2013     1     8  454    2351  
## 9  2013     1     9   2.00  2252  
## 10 2013     1    10   3.00  2320  
## # ... with 355 more rows
```

■ Measures of position: `first(x)`, `nth(x, 2)`, `last(x)`.

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(first_dep = first(dep_time), last_dep = last(dep_time))
```

```
## # A tibble: 365 x 5  
## # Groups:   year, month [?]  
##   year month   day first_dep last_dep  
##   <int> <int> <int>     <int>     <int>  
## 1  2013     1     1       517      2356  
## 2  2013     1     2        42      2354  
## 3  2013     1     3        32      2349  
## 4  2013     1     4        25      2358  
## 5  2013     1     5        14      2357  
## 6  2013     1     6        16      2355  
## 7  2013     1     7        49      2359  
## 8  2013     1     8       454      2351  
## 9  2013     1     9         2      2252  
## 10 2013     1    10         3      2320  
## # ... with 355 more rows
```

- Counts: `n(x)`, `sum(!is.na(x))`, `n_distinct(x)`.

```
not_cancelled %>%  
  group_by(dest) %>%  
  summarize(carriers = n_distinct(carrier)) %>%  
  arrange(desc(carriers))
```

```
## # A tibble: 104 x 2  
##   dest carriers  
##   <chr>    <int>  
## 1 ATL      7  
## 2 BOS      7  
## 3 CLT      7  
## 4 ORD      7  
## 5 TPA      7  
## 6 AUS      6  
## 7 DCA      6  
## 8 DTW      6  
## 9 IAD      6  
## 10 MSP     6  
## # ... with 94 more rows
```

- A simple helper function for counts:

```
not_cancelled %>% count(dest)
```

```
## # A tibble: 104 x 2
##   dest      n
##   <chr> <int>
## 1 ABQ     254
## 2 ACK     264
## 3 ALB     418
## 4 ANC       8
## 5 ATL   16837
## 6 AUS    2411
## 7 AVL     261
## 8 BDL     412
## 9 BGR     358
## 10 BHM    269
## # ... with 94 more rows
```

■ Counts with an optional weight variable:

```
not_cancelled %>% count(tailnum, wt = distance)
```

```
## # A tibble: 4,037 x 2
##   tailnum      n
##   <chr>    <dbl>
## 1 D942DN    3418
## 2 NOEGMQ 239143
## 3 N10156 109664
## 4 N102UW  25722
## 5 N103US  24619
## 6 N104UW  24616
## 7 N10575 139903
## 8 N105UW  23618
## 9 N107US  21677
## 10 N108UW  32070
## # ... with 4,027 more rows
```


- Counts of logical values: e.g., `sum(x > 10)`.

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(n_early = sum(dep_time < 500))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day n_early  
##   <int> <int> <int>   <int>  
## 1  2013     1     1       0  
## 2  2013     1     2       3  
## 3  2013     1     3       4  
## 4  2013     1     4       3  
## 5  2013     1     5       3  
## 6  2013     1     6       2  
## 7  2013     1     7       2  
## 8  2013     1     8       1  
## 9  2013     1     9       3  
## 10 2013     1    10       3  
## # ... with 355 more rows
```

- Proportions of logical values: e.g., `mean(y == 0)`.

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(hour_perc = mean(arr_delay > 60))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day hour_perc  
##   <int> <int> <int>     <dbl>  
## 1  2013     1     1    0.0722  
## 2  2013     1     2    0.0851  
## 3  2013     1     3    0.0567  
## 4  2013     1     4    0.0396  
## 5  2013     1     5    0.0349  
## 6  2013     1     6    0.0470  
## 7  2013     1     7    0.0333  
## 8  2013     1     8    0.0213  
## 9  2013     1     9    0.0202  
## 10 2013     1    10    0.0183  
## # ... with 355 more rows
```

Grouping by multiple variables I

```
daily <- group_by(flights, year, month, day)
(per_day <- summarize(daily, flights = n()))
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [?]
##   year month   day flights
##   <int> <int> <int>    <int>
## 1  2013     1     1     842
## 2  2013     1     2     943
## 3  2013     1     3     914
## 4  2013     1     4     915
## 5  2013     1     5     720
## 6  2013     1     6     832
## 7  2013     1     7     933
## 8  2013     1     8     899
## 9  2013     1     9     902
## 10 2013     1    10     932
## # ... with 355 more rows
```

Grouping by multiple variables II

```
(per_month <- summarize(per_day, flights = sum(flights)))  
(per_year <- summarize(per_month, flights = sum(flights)))
```

```
## # A tibble: 12 x 3  
## # Groups:   year [?]  
##   year month flights  
##   <int> <int>   <int>  
## 1  2013     1  27004  
## 2  2013     2  24951  
## 3  2013     3  28834  
## 4  2013     4  28330  
## 5  2013     5  28796  
## 6  2013     6  28243  
## 7  2013     7  29425  
## 8  2013     8  29327  
## 9  2013     9  27574  
## 10 2013    10  28889  
## 11 2013    11  27268  
## 12 2013    12  28135  
## # A tibble: 1 x 2  
##   year flights  
##   <int>   <int>  
## 1  2013  336776
```

```
daily %>%  
  ungroup() %>%           # no longer grouped by date  
  summarize(flights = n()) # all flights
```

```
## # A tibble: 1 x 1  
##   flights  
##   <int>  
## 1  336776
```

```
(popular_dests <- flights %>%  
  group_by(dest) %>%  
  filter(n() > 365))
```

```
## # A tibble: 332,577 x 19  
## # Groups:   dest [77]  
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>  
## 1  2013     1     1     517           515        2.00     830  
## 2  2013     1     1     533           529        4.00     850  
## 3  2013     1     1     542           540        2.00     923  
## 4  2013     1     1     544           545       -1.00    1004  
## 5  2013     1     1     554           600       -6.00     812  
## 6  2013     1     1     554           558       -4.00     740  
## 7  2013     1     1     555           600       -5.00     913  
## 8  2013     1     1     557           600       -3.00     709  
## 9  2013     1     1     557           600       -3.00     838  
## 10 2013     1     1     558           600       -2.00     753  
## # ... with 332,567 more rows, and 12 more variables:  
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,  
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
## #   time_hour <dtm>
```

```
popular_dests %>%  
  filter(arr_delay > 0) %>%  
  mutate(prop_delay = arr_delay / sum(arr_delay)) %>%  
  select(year:day, dest, arr_delay, prop_delay)
```

```
## # A tibble: 131,106 x 6  
## # Groups:   dest [77]  
##   year month   day dest  arr_delay prop_delay  
##   <int> <int> <int> <chr>    <dbl>      <dbl>  
## 1  2013     1     1 IAH      11.0    0.000111  
## 2  2013     1     1 IAH      20.0    0.000201  
## 3  2013     1     1 MIA      33.0    0.000235  
## 4  2013     1     1 ORD      12.0    0.0000424  
## 5  2013     1     1 FLL      19.0    0.0000938  
## 6  2013     1     1 ORD       8.00    0.0000283  
## 7  2013     1     1 LAX       7.00    0.0000344  
## 8  2013     1     1 DFW      31.0    0.000282  
## 9  2013     1     1 ATL      12.0    0.0000400  
## 10 2013     1     1 DTW      16.0    0.000116  
## # ... with 131,096 more rows
```

- 1 Overview
- 2 Tibbles
- 3 Data manipulation
- 4 The `%>%` operator
- 5 More on data manipulation
- 6 Tidy data**

“Happy families are all alike; every unhappy family is unhappy in its own way.” — Leo Tolstoy

“Tidy datasets are all alike, but every messy dataset is messy in its own way.” — Hadley Wickham

To learn more about the underlying theory, see the [Tidy Data paper](#).

```
table1
```

```
## # A tibble: 6 x 4
##   country      year cases population
##   <chr>      <int> <int>      <int>
## 1 Afghanistan 1999     745   19987071
## 2 Afghanistan 2000    2666   20595360
## 3 Brazil      1999   37737   172006362
## 4 Brazil      2000   80488   174504898
## 5 China       1999  212258  1272915272
## 6 China       2000  213766  1280428583
```

```
table2
```

```
## # A tibble: 12 x 4
##   country      year type      count
##   <chr>      <int> <chr>    <int>
## 1 Afghanistan 1999 cases      745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000 cases     2666
## 4 Afghanistan 2000 population 20595360
## 5 Brazil      1999 cases     37737
## 6 Brazil      1999 population 172006362
## 7 Brazil      2000 cases     80488
## 8 Brazil      2000 population 174504898
## 9 China       1999 cases     212258
## 10 China      1999 population 1272915272
## 11 China      2000 cases     213766
## 12 China      2000 population 1280428583
```

```
table3
```

```
## # A tibble: 6 x 3
##   country      year rate
## * <chr>      <int> <chr>
## 1 Afghanistan 1999 745/19987071
## 2 Afghanistan 2000 2666/20595360
## 3 Brazil      1999 37737/172006362
## 4 Brazil      2000 80488/174504898
## 5 China       1999 212258/1272915272
## 6 China       2000 213766/1280428583
```

Fourth representation

```
table4a  # cases
```

```
## # A tibble: 3 x 3
##   country      '1999' '2000'
## * <chr>      <int>  <int>
## 1 Afghanistan    745    2666
## 2 Brazil         37737  80488
## 3 China          212258 213766
```

```
table4b  # population
```

```
## # A tibble: 3 x 3
##   country      '1999'      '2000'
## * <chr>      <int>      <int>
## 1 Afghanistan 19987071  20595360
## 2 Brazil      172006362 174504898
## 3 China       1272915272 1280428583
```

What makes a dataset tidy?

Three interrelated rules:

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.

country	year	cases	population
Afghanistan		5	19907071
Afghanistan	2000	666	20005360
Brazil	1999	3737	17206362
Brazil	2000	8488	17404898
China	1999	21258	127215272
China	2000	21266	128008583

variables

country	year	cases	population
Afghanistan	2000	666	20005360
Afghanistan	1999	5	19907071
Brazil	2000	8488	17404898
Brazil	1999	3737	17206362
China	2000	21266	128008583
China	1999	21258	127215272

observations

country	year	cases	population
Afghanistan	2000	666	20005360
Afghanistan	1999	5	19907071
Brazil	2000	8488	17404898
Brazil	1999	3737	17206362
China	2000	21266	128008583
China	1999	21258	127215272

values

Because it's impossible to only satisfy two of the three:

1. Put each dataset in a tibble.
2. Put each variable in a column.

1. If you have a consistent data structure, it's easier to learn the tools that work with it because they have an underlying uniformity.
2. There's a specific advantage to placing variables in columns because it allows R's vectorized nature to shine.

The principles of tidy data seem obvious, BUT:

1. Most people aren't familiar with the principles of tidy data.
2. Data is often organised to facilitate some use other than analysis.

Hence, for most real analyses, you'll need to do some tidying.

1. Figure out what the variables and observations are.
2. Resolve one of two common problems:
 1. One variable might be spread across multiple columns.
 2. One observation might be scattered across multiple rows.

To fix these problems, you'll need `gather()` and `spread()`.

Gathering with `gather()`

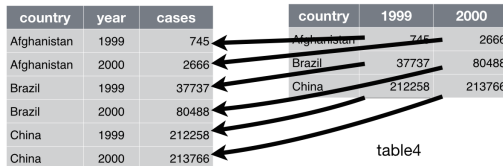
```
table4a
```

```
## # A tibble: 3 x 3
##   country      '1999' '2000'
## * <chr>      <int> <int>
## 1 Afghanistan    745   2666
## 2 Brazil        37737  80488
## 3 China         212258 213766
```

```
table4a %>% gather('1999', '2000', key = "year", value = "cases")
```

```
## # A tibble: 6 x 3
##   country      year  cases
##   <chr>      <chr> <int>
## 1 Afghanistan 1999    745
## 2 Brazil      1999  37737
## 3 China       1999 212258
## 4 Afghanistan 2000    2666
## 5 Brazil      2000  80488
## 6 China       2000 213766
```

Visual interpretation of `gather()`



country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4

Spreading with spread() I

```
table2
```

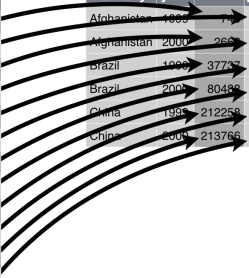
```
## # A tibble: 12 x 4
##   country      year type      count
##   <chr>      <int> <chr>    <int>
## 1 Afghanistan 1999 cases      745
## 2 Afghanistan 1999 population 19987071
## 3 Afghanistan 2000 cases      2666
## 4 Afghanistan 2000 population 20595360
## 5 Brazil      1999 cases      37737
## 6 Brazil      1999 population 172006362
## 7 Brazil      2000 cases      80488
## 8 Brazil      2000 population 174504898
## 9 China       1999 cases      212258
## 10 China      1999 population 1272915272
## 11 China      2000 cases      213766
## 12 China      2000 population 1280428583
```

Spreading with spread() II

```
table2 %>% spread(key = type, value = count)
```

```
## # A tibble: 6 x 4
##   country    year cases population
##   <chr>      <int> <int>      <int>
## 1 Afghanistan 1999    745   19987071
## 2 Afghanistan 2000   2666   20595360
## 3 Brazil      1999  37737   172006362
## 4 Brazil      2000  80488   174504898
## 5 China       1999 212258  1272915272
## 6 China       2000 213766  1280428583
```

Visual interpretation of spread()



country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

table2

```
table3
```

```
## # A tibble: 6 x 3
##   country      year rate
## * <chr>      <int> <chr>
## 1 Afghanistan  1999 745/19987071
## 2 Afghanistan  2000 2666/20595360
## 3 Brazil       1999 37737/172006362
## 4 Brazil       2000 80488/174504898
## 5 China        1999 212258/1272915272
## 6 China        2000 213766/1280428583
```

```
table3 %>% separate(rate, into = c("cases", "population"))
```

```
## # A tibble: 6 x 4
##   country    year cases population
## * <chr>      <int> <chr>   <chr>
## 1 Afghanistan 1999  745    19987071
## 2 Afghanistan 2000 2666    20595360
## 3 Brazil       1999 37737   172006362
## 4 Brazil       2000 80488   174504898
## 5 China        1999 212258  1272915272
## 6 China        2000 213766  1280428583
```

Visual interpretation of `separate()`



The diagram illustrates the `separate()` function in R. It shows a curved arrow originating from the `rate` column of the `table3` table on the left and pointing to the `cases` and `population` columns of the resulting table on the right. The `rate` column in `table3` contains values in the format `cases / population`, which are then split into two separate columns.

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

table3

separate() using convert = TRUE

```
table3 %>%  
  separate(rate, into = c("cases", "population"), convert = TRUE)
```

```
## # A tibble: 6 x 4  
##   country      year cases population  
## * <chr>      <int> <int>      <int>  
## 1 Afghanistan 1999     745   19987071  
## 2 Afghanistan 2000    2666   20595360  
## 3 Brazil       1999   37737   172006362  
## 4 Brazil       2000   80488   174504898  
## 5 China        1999  212258  1272915272  
## 6 China        2000  213766  1280428583
```

Unite two columns with unite()


```
table5
```

```
## # A tibble: 6 x 4
##   country    century year  rate
## * <chr>      <chr>   <chr> <chr>
## 1 Afghanistan 19      99    745/19987071
## 2 Afghanistan 20      00    2666/20595360
## 3 Brazil      19      99    37737/172006362
## 4 Brazil      20      00    80488/174504898
## 5 China       19      99    212258/1272915272
## 6 China       20      00    213766/1280428583
```

```
table5 %>% unite(new, century, year, sep = "")
```

```
## # A tibble: 6 x 3
##   country    new  rate
##   <chr>      <chr> <chr>
## 1 Afghanistan 1999  745/19987071
## 2 Afghanistan 2000  2666/20595360
## 3 Brazil      1999  37737/172006362
## 4 Brazil      2000  80488/174504898
## 5 China       1999  212258/1272915272
## 6 China       2000  213766/1280428583
```

Visual interpretation of unite()



A value can be missing in one of two possible ways:

- **Explicitly**, i.e. flagged with NA.
- **Implicitly**, i.e. simply not present in the data.

*“An explicit missing value is the presence of an absence;
an implicit missing value is the absence of a presence.”*
— Hadley Wickham

Are there missing values in this dataset?

```
stocks <- tibble(  
  year   = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),  
  qtr    = c( 1,    2,    3,    4,    2,    3,    4),  
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66)  
)
```

```
stocks %>%  
  spread(year, return)
```

```
## # A tibble: 4 x 3  
##   qtr '2015' '2016'  
##   <dbl> <dbl> <dbl>  
## 1  1.00  1.88  NA  
## 2  2.00  0.590  0.920  
## 3  3.00  0.350  0.170  
## 4  4.00  NA      2.66
```

```
stocks %>%  
  spread(year, return) %>%  
  gather(year, return, '2015':'2016', na.rm = TRUE)
```

```
## # A tibble: 6 x 3  
##   qtr year  return  
## * <dbl> <chr>  <dbl>  
## 1  1.00 2015    1.88  
## 2  2.00 2015    0.590  
## 3  3.00 2015    0.350  
## 4  2.00 2016    0.920  
## 5  3.00 2016    0.170  
## 6  4.00 2016    2.66
```

```
stocks %>% complete(year, qtr)
```

```
## # A tibble: 8 x 3
##   year    qtr return
##   <dbl> <dbl> <dbl>
## 1  2015    1.00  1.88
## 2  2015    2.00  0.590
## 3  2015    3.00  0.350
## 4  2015    4.00  NA
## 5  2016    1.00  NA
## 6  2016    2.00  0.920
## 7  2016    3.00  0.170
## 8  2016    4.00  2.66
```

Fill in missing values with fill()

```
treatment <- tribble(
  ~ person, ~ treatment, ~ response,
  "Derrick Whitmore", 1, 7,
  NA, 2, 10,
  NA, 3, 9,
  "Katherine Burke", 1, 4
)
treatment %>%
  fill(person)
```

```
## # A tibble: 4 x 3
##   person      treatment response
##   <chr>      <dbl>     <dbl>
## 1 Derrick Whitmore    1.00     7.00
## 2 Derrick Whitmore    2.00    10.0
## 3 Derrick Whitmore    3.00     9.00
## 4 Katherine Burke     1.00     4.00
```