

Lecture 5: Data Wrangling II

Data Science for Business Analytics

Thibault Vatter <thibault.vatter@unil.ch>

Department of Statistics, Columbia University and HEC Lausanne, UNIL

26.03.2018

1 Relational data

2 Combining tables

3 Dates and times

4 Factors

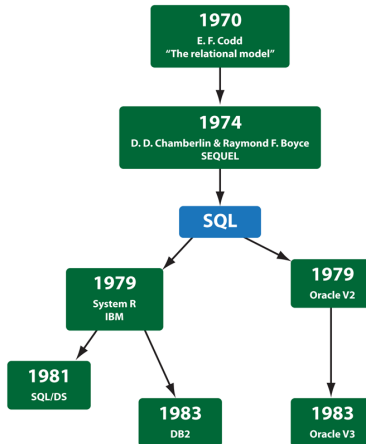
5 Strings

- Until now: analysis of a single table of data.
- Typically: multiple tables of data to be combined.

Multiple tables of data are called **relational data**:

- Because relations, not just the individual datasets, are important.
- Relations are always defined for a pair of tables.
- Relations of three or more tables are built from the relations between pairs.

- Common place to find relational data.
- Oracle, MySQL, Microsoft SQL Server, PostgreSQL, IBM DB2, Microsoft Access, SQLite, and others.



All 336,776 flights that departed from NYC in 2013 (US BTS):

```
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
##  1  2013     1     1     517           515           2.00     830
##  2  2013     1     1     533           529           4.00     850
##  3  2013     1     1     542           540           2.00     923
##  4  2013     1     1     544           545          -1.00    1004
##  5  2013     1     1     554           600          -6.00     812
##  6  2013     1     1     554           558          -4.00     740
##  7  2013     1     1     555           600          -5.00     913
##  8  2013     1     1     557           600          -3.00     709
##  9  2013     1     1     557           600          -3.00     838
## 10  2013     1     1     558           600          -2.00     753
## # ... with 336,766 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
airlines
```

```
## # A tibble: 16 x 2
##   carrier name
##   <chr>    <chr>
## 1 9E      Endeavor Air Inc.
## 2 AA      American Airlines Inc.
## 3 AS      Alaska Airlines Inc.
## 4 B6      JetBlue Airways
## 5 DL      Delta Air Lines Inc.
## 6 EV      ExpressJet Airlines Inc.
## 7 F9      Frontier Airlines Inc.
## 8 FL      AirTran Airways Corporation
## 9 HA      Hawaiian Airlines Inc.
## 10 MQ     Envoy Air
## 11 OO     SkyWest Airlines Inc.
## 12 UA     United Air Lines Inc.
## 13 US     US Airways Inc.
## 14 VX     Virgin America
## 15 WN     Southwest Airlines Co.
## 16 YV     Mesa Airlines Inc.
```

airports

```
## # A tibble: 1,458 x 8
##   faa   name          lat    lon    alt    tz dst  tzone
##   <chr> <chr>         <dbl> <dbl> <int> <dbl> <chr> <chr>
## 1 04G   Lansdowne Airport 41.1 - 80.6 1044 -5.00 A   America/Ne~
## 2 06A   Moton Field Muni~ 32.5 - 85.7 264 -6.00 A   America/Ch~
## 3 06C   Schaumburg Regio~ 42.0 - 88.1 801 -6.00 A   America/Ch~
## 4 06N   Randall Airport  41.4 - 74.4 523 -5.00 A   America/Ne~
## 5 09J   Jekyll Island Ai~ 31.1 - 81.4 11 -5.00 A   America/Ne~
## 6 0A9   Elizabethton Mun~ 36.4 - 82.2 1593 -5.00 A   America/Ne~
## 7 0G6   Williams County ~ 41.5 - 84.5 730 -5.00 A   America/Ne~
## 8 0G7   Finger Lakes Reg~ 42.9 - 76.8 492 -5.00 A   America/Ne~
## 9 0P2   Shoestring Aviat~ 39.8 - 76.6 1000 -5.00 U   America/Ne~
## 10 0S9   Jefferson County~ 48.1 -123    108 -8.00 A   America/Lo~
## # ... with 1,448 more rows
```

planes

```
## # A tibble: 3,322 x 9
##   tailnum  year type  manufacturer model engines seats speed engine
##   <chr>    <int> <chr>  <chr>          <chr>    <int> <int> <int> <chr>
## 1 N10156   2004 Fixed~ EMBRAER        EMB~      2    55    NA Turbo~
## 2 N102UW   1998 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 3 N103US   1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 4 N104UW   1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 5 N10575   2002 Fixed~ EMBRAER        EMB~      2    55    NA Turbo~
## 6 N105UW   1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 7 N107US   1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 8 N108UW   1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 9 N109UW   1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## 10 N110UW  1999 Fixed~ AIRBUS INDU~   A320~    2   182    NA Turbo~
## # ... with 3,312 more rows
```

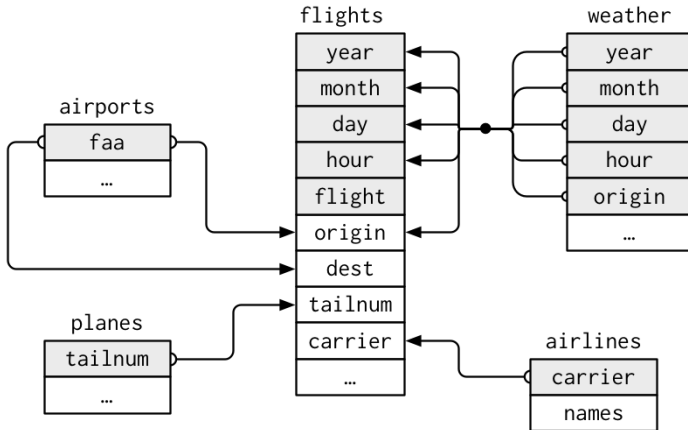


```
weather
```

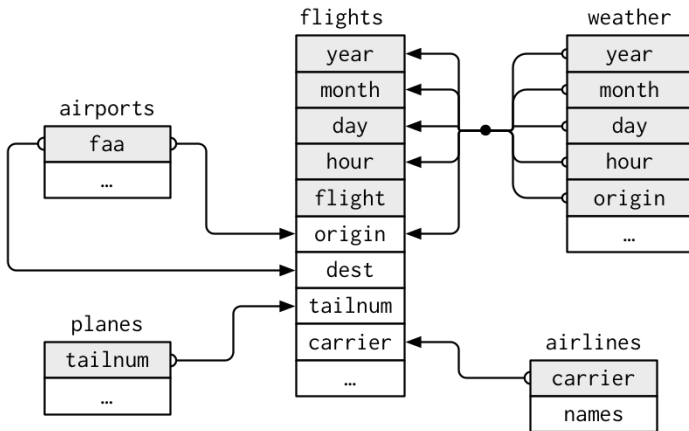
```
## # A tibble: 26,130 x 15
```

```
##   origin year month   day hour temp dewp humid wind_dir
##   <chr>  <dbl> <dbl> <int> <int> <dbl> <dbl> <dbl>    <dbl>
## 1 EWR    2013  1.00     1     0  37.0  21.9  54.0      230
## 2 EWR    2013  1.00     1     1  37.0  21.9  54.0      230
## 3 EWR    2013  1.00     1     2  37.9  21.9  52.1      230
## 4 EWR    2013  1.00     1     3  37.9  23.0  54.5      230
## 5 EWR    2013  1.00     1     4  37.9  24.1  57.0      240
## 6 EWR    2013  1.00     1     6  39.0  26.1  59.4      270
## 7 EWR    2013  1.00     1     7  39.0  27.0  61.6      250
## 8 EWR    2013  1.00     1     8  39.0  28.0  64.4      240
## 9 EWR    2013  1.00     1     9  39.9  28.0  62.2      250
## 10 EWR   2013  1.00     1    10  39.0  28.0  64.4      260
```

```
## # ... with 26,120 more rows, and 6 more variables: wind_speed <dbl>,
## #   wind_gust <dbl>, precip <dbl>, pressure <dbl>, visib <dbl>,
## #   time_hour <dttm>
```

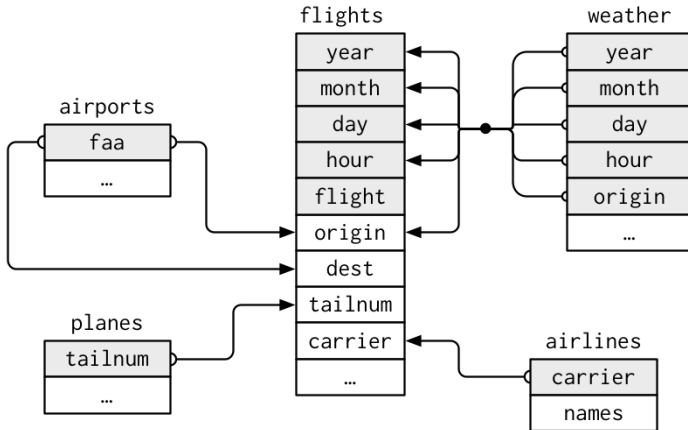


Exercise 1



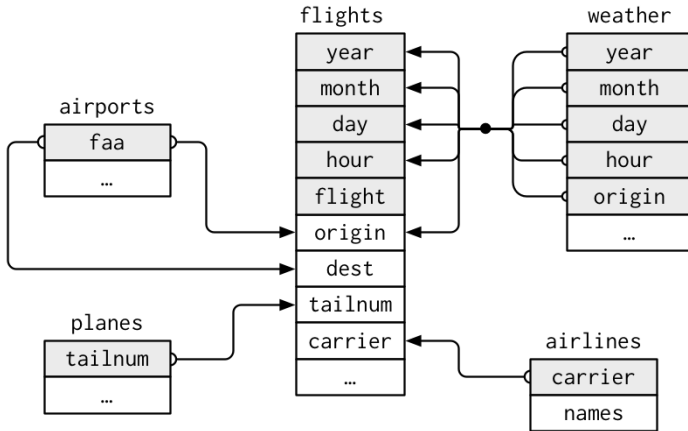
Imagine you wanted to draw (approximately) the route each plane flies from its origin to its destination. What variables would you need? What tables would you need to combine?

Exercise 2



I forgot to draw the relationship between weather and airports.
What is the relationship and how should it appear in the diagram?

Exercise 3



weather only contains information for the origin (NYC) airports. If it contained weather records for all airports in the USA, what additional relation would it define with `flights`?

- Variables used to connect pair of tables.
- Uniquely identifies an observation.
- Either a single variable (e.g., `tailnum` for planes) or multiple variables (e.g., `year`, `month`, `day`, `hour`, and `origin` for weather).

Two types of **keys**:

- A **primary key** uniquely identifies an observation **in its own table** (e.g., `planes$tailnum`).
- A **foreign key** uniquely identifies an observation **in another table** (e.g., `flights$tailnum`).

Note that:

- A variable can be both a primary key *and* a foreign key.
- A primary key and the corresponding foreign key in another table form a **relation**.
- Relations are typically one-to-many (e.g., `flights` and `planes`).

Is a given key primary?

```
planes %>%  
  count(tailnum) %>%  
  filter(n > 1)
```

```
## # A tibble: 0 x 2  
## # ... with 2 variables: tailnum <chr>, n <int>
```

```
weather %>%  
  count(year, month, day, hour, origin) %>%  
  filter(n > 1)
```

```
## # A tibble: 0 x 6  
## # ... with 6 variables: year <dbl>, month <dbl>, day <int>,  
## #   hour <int>, origin <chr>, n <int>
```

No explicit primary key?

```
flights %>%  
  count(year, month, day, flight) %>%  
  filter(n > 1)
```

```
## # A tibble: 29,768 x 5  
##   year month   day flight     n  
##   <int> <int> <int>   <int> <int>  
## 1  2013     1     1       1     2  
## 2  2013     1     1       3     2  
## 3  2013     1     1       4     2  
## 4  2013     1     1      11     3  
## 5  2013     1     1      15     2  
## 6  2013     1     1      21     2  
## 7  2013     1     1      27     4  
## 8  2013     1     1      31     2  
## 9  2013     1     1      32     2  
## 10 2013     1     1      35     2  
## # ... with 29,758 more rows
```

- Solution: add one with `mutate()` and `row_number()`.
- This is called a **surrogate key**.

1 Relational data

2 Combining tables

3 Dates and times

4 Factors

5 Strings

Three families of verbs to work with relational data:

- **Mutating joins**, which add new variables to one data frame from matching observations in another.
- **Filtering joins**, which filter observations from one data frame based on whether or not they match an observation in the other table.
- **Set operations**, which treat observations as if they were set elements.

Create a narrower dataset

```
flights2 <- flights %>%  
  select(year:day, hour, origin, dest, tailnum, carrier)  
flights2
```

```
## # A tibble: 336,776 x 8  
##   year month   day hour origin dest  tailnum carrier  
##   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr>  
## 1  2013     1     1  5.00 EWR   IAH   N14228 UA  
## 2  2013     1     1  5.00 LGA   IAH   N24211 UA  
## 3  2013     1     1  5.00 JFK   MIA   N619AA AA  
## 4  2013     1     1  5.00 JFK   BQN   N804JB B6  
## 5  2013     1     1  6.00 LGA   ATL   N668DN DL  
## 6  2013     1     1  5.00 EWR   ORD   N39463 UA  
## 7  2013     1     1  6.00 EWR   FLL   N516JB B6  
## 8  2013     1     1  6.00 LGA   IAD   N829AS EV  
## 9  2013     1     1  6.00 JFK   MCO   N593JB B6  
## 10 2013     1     1  6.00 LGA   ORD   N3ALAA AA  
## # ... with 336,766 more rows
```

A simple example

```
flights2 %>%  
  select(-origin, -dest) %>%  
  left_join(airlines, by = "carrier")
```

```
## # A tibble: 336,776 x 7
```

```
##   year month   day hour tailnum carrier name  
##   <int> <int> <int> <dbl> <chr>   <chr>   <chr>  
## 1  2013     1     1  5.00 N14228  UA      United Air Lines Inc.  
## 2  2013     1     1  5.00 N24211  UA      United Air Lines Inc.  
## 3  2013     1     1  5.00 N619AA  AA      American Airlines Inc.  
## 4  2013     1     1  5.00 N804JB  B6      JetBlue Airways  
## 5  2013     1     1  6.00 N668DN  DL      Delta Air Lines Inc.  
## 6  2013     1     1  5.00 N39463  UA      United Air Lines Inc.  
## 7  2013     1     1  6.00 N516JB  B6      JetBlue Airways  
## 8  2013     1     1  6.00 N829AS  EV      ExpressJet Airlines Inc.  
## 9  2013     1     1  6.00 N593JB  B6      JetBlue Airways  
## 10 2013     1     1  6.00 N3ALAA  AA      American Airlines Inc.  
## # ... with 336,766 more rows
```

Why mutating join?

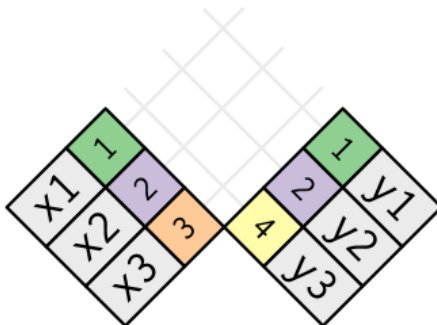
```
flights2 %>%  
  select(-origin, -dest) %>%  
  mutate(name = airlines$name[match(carrier, airlines$carrier)])
```

```
## # A tibble: 336,776 x 7
```

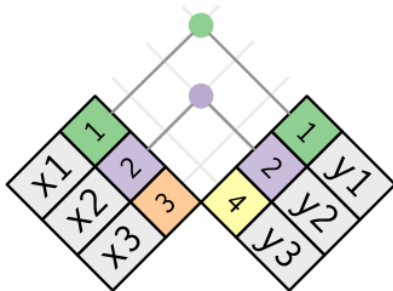
```
##   year month   day hour tailnum carrier name  
##   <int> <int> <int> <dbl> <chr>   <chr>   <chr>  
## 1  2013     1     1  5.00 N14228   UA      United Air Lines Inc.  
## 2  2013     1     1  5.00 N24211   UA      United Air Lines Inc.  
## 3  2013     1     1  5.00 N619AA   AA      American Airlines Inc.  
## 4  2013     1     1  5.00 N804JB   B6      JetBlue Airways  
## 5  2013     1     1  6.00 N668DN   DL      Delta Air Lines Inc.  
## 6  2013     1     1  5.00 N39463   UA      United Air Lines Inc.  
## 7  2013     1     1  6.00 N516JB   B6      JetBlue Airways  
## 8  2013     1     1  6.00 N829AS   EV      ExpressJet Airlines Inc.  
## 9  2013     1     1  6.00 N593JB   B6      JetBlue Airways  
## 10 2013     1     1  6.00 N3ALAA   AA      American Airlines Inc.  
## # ... with 336,766 more rows
```

Understanding mutating joins

```
x <- tribble(~key, ~val_x,  
  1, "x1",  
  2, "x2",  
  3, "x3")  
y <- tribble(~key, ~val_y,  
  1, "y1",  
  2, "y2",  
  4, "y3")
```



Inner join



key	val_x	val_y
1	x1	y1
2	x2	y2

```
x %>%  
  inner_join(y, by = "key")
```

```
## # A tibble: 2 x 3  
##   key val_x val_y  
##   <dbl> <chr> <chr>  
## 1  1.00 x1    y1  
## 2  2.00 x2    y2
```

An **outer join** keeps observations that appear in at least one of the tables:

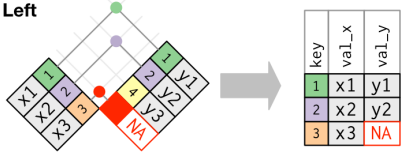
- A **left join** keeps all observations in x .
- A **right join** keeps all observations in y .
- A **full join** keeps all observations in x and y .

They work by adding to each table an additional “virtual” observation which

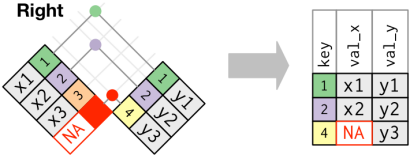
- has a key that always matches (if no other key matches),
- and a value filled with NA.

Outer joins II

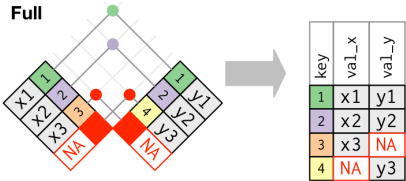
Left



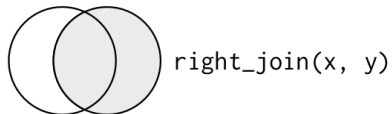
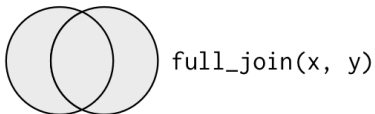
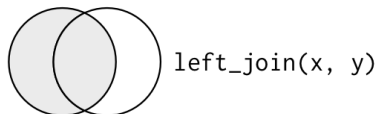
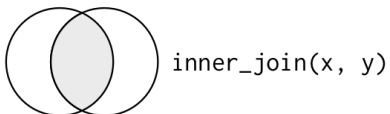
Right



Full



A Venn diagram for joins



Two possibilities:

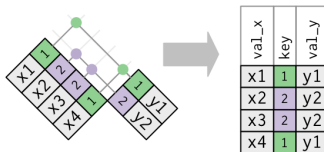
1. One table has duplicate keys.

- Useful to add in additional information as there is typically a one-to-many relationship.

2. Both tables have duplicate keys.

- Usually an error because in neither table do the keys uniquely identify an observation.
- When you join duplicated keys, you get all possible combinations (i.e., the Cartesian product).

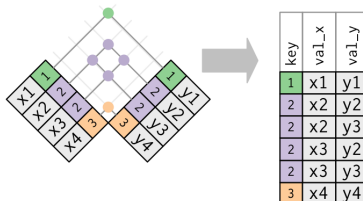
One table has duplicate keys



```
x <- tribble(~key, ~val_x,  
             1, "x1",  
             2, "x2",  
             2, "x3",  
             1, "x4")  
y <- tribble(~key, ~val_y,  
             1, "y1",  
             2, "y2")  
left_join(x, y, by = "key")
```

```
## # A tibble: 4 x 3  
##   key val_x val_y  
##   <dbl> <chr> <chr>  
## 1  1.00 x1    y1  
## 2  2.00 x2    y2  
## 3  2.00 x3    y2  
## 4  1.00 x4    y1
```

Both tables have duplicate keys



```
x <- tribble(~key, ~val_x, 1, "x1", 2, "x2", 2, "x3", 3, "x4")
y <- tribble(~key, ~val_y, 1, "y1", 2, "y2", 2, "y3", 3, "y4")
left_join(x, y, by = "key")
```

```
## # A tibble: 6 x 3
##   key val_x val_y
##   <dbl> <chr> <chr>
## 1  1.00 x1    y1
## 2  2.00 x2    y2
## 3  2.00 x2    y3
## 4  2.00 x3    y2
## 5  2.00 x3    y3
## 6  3.00 x4    y4
```

Defining the key columns

Default uses all variables that appear in both tables (**natural**):

```
flights2 %>%  
  left_join(weather)
```

```
## Joining, by = c("year", "month", "day", "hour", "origin")  
  
## # A tibble: 336,776 x 18  
##   year month   day hour origin dest tailnum carrier temp dewp  
##   <dbl> <dbl> <int> <dbl> <chr>  <chr> <chr>   <chr>   <dbl> <dbl>  
## 1  2013   1.00     1  5.00 EWR   IAH   N14228 UA        NA    NA  
## 2  2013   1.00     1  5.00 LGA   IAH   N24211 UA        NA    NA  
## 3  2013   1.00     1  5.00 JFK   MIA   N619AA AA        NA    NA  
## 4  2013   1.00     1  5.00 JFK   BQN   N804JB B6        NA    NA  
## 5  2013   1.00     1  6.00 LGA   ATL   N668DN DL       39.9  26.1  
## 6  2013   1.00     1  5.00 EWR   ORD   N39463 UA        NA    NA  
## 7  2013   1.00     1  6.00 EWR   FLL   N516JB B6       39.0  26.1  
## 8  2013   1.00     1  6.00 LGA   IAD   N829AS EV       39.9  26.1  
## 9  2013   1.00     1  6.00 JFK   MCO   N593JB B6       39.0  26.1  
## 10 2013   1.00     1  6.00 LGA   ORD   N3ALAA AA       39.9  26.1  
## # ... with 336,766 more rows, and 8 more variables: humid <dbl>,  
## #   wind_dir <dbl>, wind_speed <dbl>, wind_gust <dbl>, precip <dbl>,  
## #   pressure <dbl>, visib <dbl>, time_hour <dtm>
```

Using a character vector

Like a natural join, but uses only some of the common variables:

```
flights2 %>%  
  left_join(planes, by = "tailnum")
```

```
## # A tibble: 336,776 x 16  
##   year.x month   day hour origin dest tailnum carrier year.y type  
##   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <int> <chr>  
## 1  2013     1     1  5.00 EWR   IAH   N14228 UA      1999 Fixe~  
## 2  2013     1     1  5.00 LGA   IAH   N24211 UA      1998 Fixe~  
## 3  2013     1     1  5.00 JFK   MIA   N619AA AA      1990 Fixe~  
## 4  2013     1     1  5.00 JFK   BQN   N804JB B6      2012 Fixe~  
## 5  2013     1     1  6.00 LGA   ATL   N668DN DL      1991 Fixe~  
## 6  2013     1     1  5.00 EWR   ORD   N39463 UA      2012 Fixe~  
## 7  2013     1     1  6.00 EWR   FLL   N516JB B6      2000 Fixe~  
## 8  2013     1     1  6.00 LGA   IAD   N829AS EV      1998 Fixe~  
## 9  2013     1     1  6.00 JFK   MCO   N593JB B6      2004 Fixe~  
## 10 2013     1     1  6.00 LGA   ORD   N3ALAA AA      NA <NA>  
## # ... with 336,766 more rows, and 6 more variables:  
## #   manufacturer <chr>, model <chr>, engines <int>, seats <int>,  
## #   speed <int>, engine <chr>
```

Using a named character vector

With `by = c("a" = "b")`, `left_join` matches variable `a` in table `x` to variable `b` in table `y`:

```
flights2 %>%  
  left_join(airports, c("dest" = "faa"))
```

```
## # A tibble: 336,776 x 15  
##   year month   day hour origin dest  tailnum carrier name    lat  
##   <int> <int> <int> <dbl> <chr> <chr> <chr>   <chr>   <chr>   <dbl>  
## 1  2013     1     1  5.00 EWR   IAH   N14228  UA      George~ 30.0  
## 2  2013     1     1  5.00 LGA   IAH   N24211  UA      George~ 30.0  
## 3  2013     1     1  5.00 JFK   MIA   N619AA  AA      Miami ~ 25.8  
## 4  2013     1     1  5.00 JFK   BQN   N804JB  B6      <NA>    NA  
## 5  2013     1     1  6.00 LGA   ATL   N668DN  DL      Hartsf~ 33.6  
## 6  2013     1     1  5.00 EWR   ORD   N39463  UA      Chicag~ 42.0  
## 7  2013     1     1  6.00 EWR   FLL   N516JB  B6      Fort L~ 26.1  
## 8  2013     1     1  6.00 LGA   IAD   N829AS  EV      Washin~ 38.9  
## 9  2013     1     1  6.00 JFK   MCO   N593JB  B6      Orland~ 28.4  
## 10 2013     1     1  6.00 LGA   ORD   N3ALAA  AA      Chicag~ 42.0  
## # ... with 336,766 more rows, and 5 more variables: lon <dbl>,  
## #   alt <int>, tz <dbl>, dst <chr>, tzone <chr>
```


`base::merge()` can perform all four types of mutating join:

dplyr	merge
<code>inner_join(x, y)</code>	<code>merge(x, y)</code>
<code>left_join(x, y)</code>	<code>merge(x, y, all.x = TRUE)</code>
<code>right_join(x, y)</code>	<code>merge(x, y, all.y = TRUE),</code>
<code>full_join(x, y)</code>	<code>merge(x, y, all.x = TRUE, all.y = TRUE)</code>

Advantages of the specific dplyr verbs:

- More clearly convey the intent of your code.
- Considerably faster and don't mess with the order of the rows.

SQL is the inspiration for dplyr's conventions:

dplyr	SQL
<code>inner_join(x, y, by = "z")</code>	<code>SELECT * FROM x INNER JOIN y USING (z)</code>
<code>left_join(x, y, by = "z")</code>	<code>SELECT * FROM x LEFT OUTER JOIN y USING (z)</code>
<code>right_join(x, y, by = "z")</code>	<code>SELECT * FROM x RIGHT OUTER JOIN y USING (z)</code>
<code>full_join(x, y, by = "z")</code>	<code>SELECT * FROM x FULL OUTER JOIN y USING (z)</code>

Note that:

- “INNER” and “OUTER” are optional, and often omitted.
- Joining different variables between the tables, e.g.
`inner_join(x, y, by = c("a" = "b"))` uses a slightly different syntax in SQL: `SELECT * FROM x INNER JOIN y ON x.a = y.b.`

Similar to mutating joins, but affect the observations rather than the variables:

- `semi_join(x, y)` **keeps** all observations in `x` that have a match in `y`.
 - ▶ Useful for matching filtered summary tables back to the original rows.
- `anti_join(x, y)` **drops** all observations in `x` that have a match in `y`.
 - ▶ Useful for diagnosing join mismatches.

Flights that went to top destinations COLUMBIA UNIVERSITY IN THE CITY OF NEW YORK

```
top_dest <- flights %>% count(dest, sort = TRUE) %>% head(10)
flights %>% filter(dest %in% top_dest$dest)
```

```
## # A tibble: 141,145 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     542             540           2.00     923
## 2  2013     1     1     554             600          -6.00     812
## 3  2013     1     1     554             558          -4.00     740
## 4  2013     1     1     555             600          -5.00     913
## 5  2013     1     1     557             600          -3.00     838
## 6  2013     1     1     558             600          -2.00     753
## 7  2013     1     1     558             600          -2.00     924
## 8  2013     1     1     558             600          -2.00     923
## 9  2013     1     1     559             559           0         702
## 10 2013     1     1     600             600           0         851
## # ... with 141,135 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

But it's difficult to extend that approach to multiple variables.

Only keeps rows in x having a match in y:

```
flights %>% semi_join(top_dest)
```

```
## Joining, by = "dest"
```

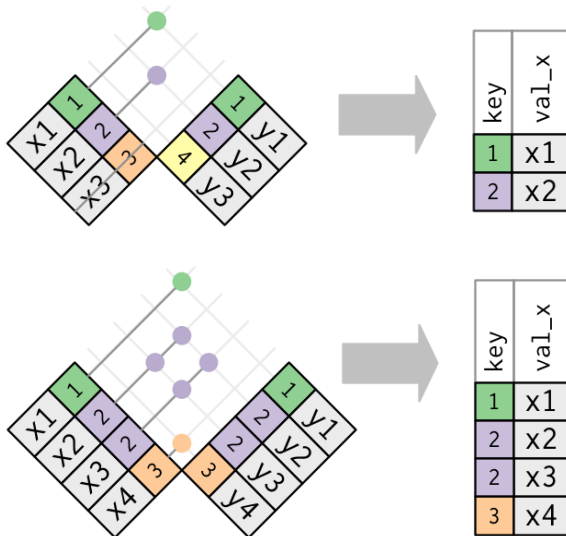
```
## # A tibble: 141,145 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>
## 1  2013     1     1     542             540          2.00     923
## 2  2013     1     1     554             600         -6.00     812
## 3  2013     1     1     554             558         -4.00     740
## 4  2013     1     1     555             600         -5.00     913
## 5  2013     1     1     557             600         -3.00     838
## 6  2013     1     1     558             600         -2.00     753
## 7  2013     1     1     558             600         -2.00     924
## 8  2013     1     1     558             600         -2.00     923
## 9  2013     1     1     559             559           0       702
##10  2013     1     1     600             600           0       851
```

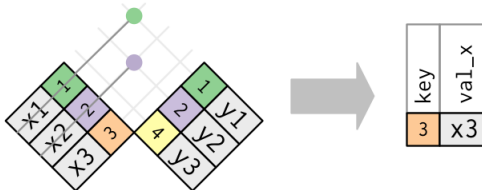
```
## # ... with 141,135 more rows, and 12 more variables:
```

```
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

Visually understand the semi-join



flights without match in planes



```
flights %>% anti_join(planes, by = "tailnum") %>% count(tailnum, sort = TRUE)
```

```
## # A tibble: 722 x 2
```

```
##   tailnum      n
```

```
##   <chr>    <int>
```

```
## 1 <NA>    2512
```

```
## 2 N725MQ   575
```

```
## 3 N722MQ   513
```

```
## 4 N723MQ   507
```

```
## 5 N713MQ   483
```

```
## 6 N735MQ   396
```

```
## 7 NOEGMQ   371
```

```
## 8 N534MQ   364
```

```
## 9 N542MQ   363
```

- Used the least frequently
- Work with a complete row, comparing the values of every variable.
- Expect the x and y inputs to have the same variables, and treat the observations like sets.

The three set operations:

- `intersect(x, y)`: return only observations in both x and y .
- `union(x, y)`: return unique observations in x and y .
- `setdiff(x, y)`: return observations in x , but not in y .

Intersect and union

```
df1 <- tribble(~x, ~y,  
              1, 1,  
              2, 1)  
df2 <- tribble(~x, ~y,  
              1, 1,  
              1, 2)
```

```
intersect(df1, df2)
```

```
## # A tibble: 1 x 2  
##       x       y  
##   <dbl> <dbl>  
## 1  1.00  1.00
```

```
union(df1, df2)
```

```
## # A tibble: 3 x 2  
##       x       y  
##   <dbl> <dbl>  
## 1  1.00  2.00  
## 2  2.00  1.00  
## 3  1.00  1.00
```

```
df1 <- tribble(~x, ~y,  
              1,  1,  
              2,  1)  
df2 <- tribble(~x, ~y,  
              1,  1,  
              1,  2)
```

```
setdiff(df1, df2)
```

```
## # A tibble: 1 x 2  
##       x     y  
##   <dbl> <dbl>  
## 1  2.00  1.00
```

```
setdiff(df2, df1)
```

```
## # A tibble: 1 x 2  
##       x     y  
##   <dbl> <dbl>  
## 1  1.00  2.00
```

1 Relational data

2 Combining tables

3 Dates and times

4 Factors

5 Strings

- Does every year have 365 days?
- Does every day have 24 hours?
- Does every minute have 60 seconds?

Three types of date/time data:

- A **date**. Tibbles print this as `<date>`.
- A **time** within a day. Tibbles print this as `<time>`.
- A **date-time** is a date plus a time: it uniquely identifies an instant in time (typically to the nearest second). Tibbles print this as `<dtm>`. Elsewhere in R these are called POSIXct.

In R:

- Focus on dates/date-times because no “native” class for times.
- If you need one, look at the **hms** package.

Use the simplest possible data type satisfying your needs!

The **lubridate** package:

- Makes it easier to work with dates and times in R,
- is not part of core tidyverse because you only need it when you're working with dates/times.

```
library(lubridate)
today()
now()
```

```
## [1] "2018-03-25"
## [1] "2018-03-25 22:24:29 CEST"
```

Three other (usual) ways to create a date/time:

- From a string.
- From individual date-time components.
- From an existing date/time object (i.e., with `as_datetime(today())` or conversely `as_date(now())`).

```
ymd("2017-01-31")  
mdy("January 31st, 2017")  
dmy("31-Jan-2017")  
  
ymd_hms("2017-01-31 20:11:59")  
mdy_hm("01/31/2017 08:01")
```

```
## [1] "2017-01-31"  
## [1] "2017-01-31"  
## [1] "2017-01-31"  
## [1] "2017-01-31 20:11:59 UTC"  
## [1] "2017-01-31 08:01:00 UTC"
```

Additionally:

```
ymd(20170131)  
ymd(20170131, tz = "UTC")
```

```
## [1] "2017-01-31"  
## [1] "2017-01-31 UTC"
```

From individual components

```
flights %>%  
  select(year, month, day, hour, minute, dep_time) %>%  
  mutate(departure = make_datetime(year, month, day, hour, minute))
```

```
## # A tibble: 336,776 x 7  
##   year month   day hour minute dep_time departure  
##   <int> <int> <int> <dbl> <dbl>   <int> <dtm>  
## 1  2013     1     1  5.00  15.0     517 2013-01-01 05:15:00  
## 2  2013     1     1  5.00  29.0     533 2013-01-01 05:29:00  
## 3  2013     1     1  5.00  40.0     542 2013-01-01 05:40:00  
## 4  2013     1     1  5.00  45.0     544 2013-01-01 05:45:00  
## 5  2013     1     1  6.00   0.0     554 2013-01-01 06:00:00  
## 6  2013     1     1  5.00  58.0     554 2013-01-01 05:58:00  
## 7  2013     1     1  6.00   0.0     555 2013-01-01 06:00:00  
## 8  2013     1     1  6.00   0.0     557 2013-01-01 06:00:00  
## 9  2013     1     1  6.00   0.0     557 2013-01-01 06:00:00  
## 10 2013     1     1  6.00   0.0     558 2013-01-01 06:00:00  
## # ... with 336,766 more rows
```

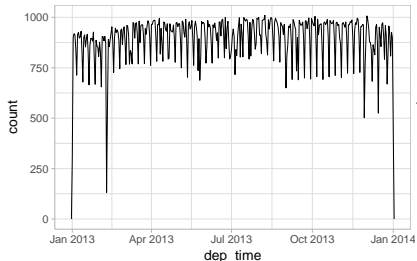

For dep_time and others such as arr_time:

```
make_datetime_100 <- function(year, month, day, time) {  
  make_datetime(year, month, day, time %% 100, time %% 100)}
```

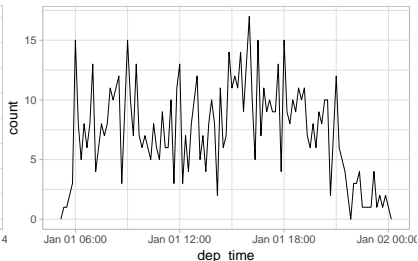
From individual components II

```
flights_dt <- flights %>% filter(!is.na(dep_time), !is.na(arr_time)) %>%  
  mutate(dep_time = make_datetime_100(year, month, day, dep_time),  
         arr_time = make_datetime_100(year, month, day, arr_time)) %>%  
  select(origin, dest, ends_with("delay"), ends_with("time"))  
  
flights_dt %>% ggplot(aes(dep_time)) +  
  geom_freqpoly(binwidth = 86400) + # 86400s = 1d  
  ggtitle("Distribution of departures in a year")  
flights_dt %>% filter(dep_time < ymd(20130102)) %>%  
  ggplot(aes(dep_time)) + geom_freqpoly(binwidth = 600) + # 600s = 10mn  
  ggtitle("Distribution of departures on January 1st")
```

Distribution of departures in a year



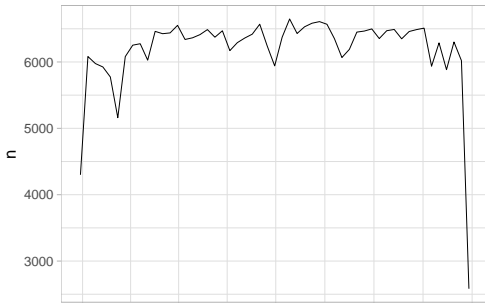
Distribution of departures on January 1st



- `floor_date()` rounds down
- `round_date()` rounds to
- `ceiling_date()` rounds up

Takes a vector of dates to adjust and then the name of the unit:

```
flights_dt %>%  
  count(week = floor_date(dep_time, "week")) %>%  
  ggplot(aes(week, n)) +  
    geom_line()
```



Getting/setting components of a date/time

Getting the components:

```
datetime <- ymd_hms("2016-07-08 12:34:56")  
c(year(datetime), month(datetime), mday(datetime),  
  yday(datetime), wday(datetime))
```

```
## [1] 2016    7    8  190    6
```

Setting the components:

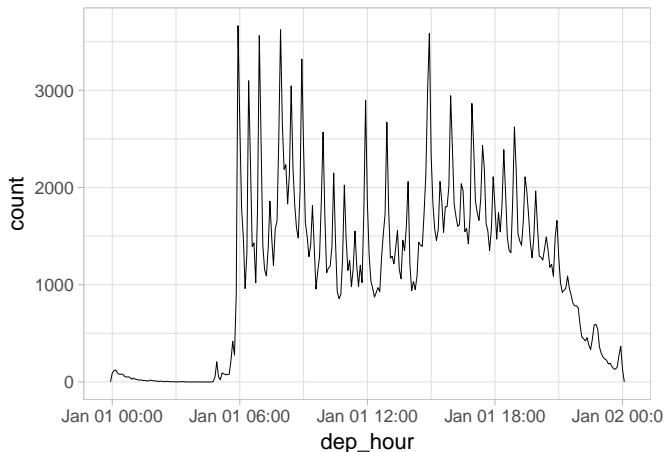
```
year(datetime) <- 2020  
datetime  
month(datetime) <- 01  
datetime  
hour(datetime) <- hour(datetime) + 1  
datetime
```

```
## [1] "2020-07-08 12:34:56 UTC"  
## [1] "2020-01-08 12:34:56 UTC"  
## [1] "2020-01-08 13:34:56 UTC"
```

Alternatively, use e.g. `update(datetime, year = 2020)`.

Flights distribution over the day

```
flights_dt %>%  
  mutate(dep_hour = update(dep_time, yday = 1)) %>%  
  ggplot(aes(dep_hour)) +  
    geom_freqpoly(binwidth = 300)
```



Goal: to do arithmetic (i.e., subtraction, addition, and division) with dates/times.

Three classes that represent time spans:

- **durations** (number of seconds).
- **periods** (human units like weeks and months).
- **intervals** (a starting and ending point).

- A **duration** always record a time span in seconds.
- Larger units created at the standard rate (60s/mn, 60mn/h, 24h/d, 7d/w, 365d/y)

```
dseconds(15)
dminutes(10)
dhours(c(12, 24))
ddays(0:5)
dweeks(3)
dyears(1)
```

```
## [1] "15s"
## [1] "600s (~10 minutes)"
## [1] "43200s (~12 hours)" "86400s (~1 days)"
## [1] "0s" "86400s (~1 days)" "172800s (~2 days)"
## [4] "259200s (~3 days)" "345600s (~4 days)" "432000s (~5 days)"
## [1] "1814400s (~3 weeks)"
## [1] "31536000s (~52.14 weeks)"
```

You can add and multiply durations:

```
2 * dyears(1)
dyears(1) + dweeks(12) + dhours(15)
```

```
## [1] "63072000s (~2 years)"
```

```
## [1] "38847600s (~1.23 years)"
```

You can add and subtract durations to and from days:

```
tomorrow <- today() + ddays(1)
last_year <- today() - dyears(1)
```

What happens here?

```
one_pm <- ymd_hms("2016-03-12 13:00:00", tz = "America/New_York")
one_pm
one_pm + ddays(1)
```

```
## [1] "2016-03-12 13:00:00 EST"
```

```
## [1] "2016-03-13 14:00:00 EDT"
```


Work with “human” times, like days (no fixed length in secs):

```
one_pm  
one_pm + days(1)
```

```
## [1] "2016-03-12 13:00:00 EST"  
## [1] "2016-03-13 13:00:00 EDT"
```

```
seconds(15)  
minutes(10)  
hours(c(12, 24))  
days(7)  
months(1:3)  
weeks(3)  
years(1)
```

```
## [1] "15S"  
## [1] "10M 0S"  
## [1] "12H 0M 0S" "24H 0M 0S"  
## [1] "7d 0H 0M 0S"  
## [1] "1m 0d 0H 0M 0S" "2m 0d 0H 0M 0S" "3m 0d 0H 0M 0S"  
## [1] "21d 0H 0M 0S"  
## [1] "1y 0m 0d 0H 0M 0S"
```

Add and multiply periods:

```
10 * (months(6) + days(1))  
days(50) + hours(25) + minutes(2)
```

```
## [1] "60m 10d 0H 0M 0S"  
## [1] "50d 25H 2M 0S"
```

Add periods to dates:

```
# A leap year  
ymd("2016-01-01") + dyears(1)  
ymd("2016-01-01") + years(1)
```

```
# Daylight Savings Time  
one_pm + ddays(1)  
one_pm + days(1)
```

```
## [1] "2016-12-31"  
## [1] "2017-01-01"  
## [1] "2016-03-13 14:00:00 EDT"  
## [1] "2016-03-13 13:00:00 EDT"
```

- What should `dyears(1) / ddays(365)` return ?
- What should `years(1) / days(1)` return ?

- What should `dyears(1) / ddays(365)` return ?
- What should `years(1) / days(1)` return ?

```
years(1) / days(1)
```

```
## estimate only: convert to intervals for accuracy
```

```
## [1] 365.25
```

The **interval** (i.e., a duration with a starting point):

```
next_year <- today() + years(1)
(today() %--% next_year) / ddays(1)
```

```
## [1] 365
```

How many periods fall into an interval:

```
(today() %--% next_year) %/% days(1)
```

```
## Note: method with signature 'Timespan#Timespan' chosen for function '%/%',
## target signature 'Interval#Period'.
## "Interval#ANY", "ANY#Period" would also be valid
```

```
## [1] 365
```

	date				date time				duration				period				interval				number			
date	-								-	+			-	+						-	+			
date time					-				-	+			-	+						-	+			
duration	-	+			-	+			-	+	/								-	+	×	/		
period	-	+			-	+						-	+						-	+	×	/		
interval											/				/									
number	-	+			-	+			-	+	×		-	+	×		-	+	×		-	+	×	/

Pick the simplest data structure that solves your problem:

- If you only care about physical time, use a duration.
- If you need to add human times, use a period.
- If you need to figure out how long a span is in human units, use an interval.

```
Sys.timezone()
```

```
## [1] "Europe/Paris"
```

```
length(OlsonNames())
```

```
## [1] 592
```

```
head(OlsonNames())
```

```
## [1] "Africa/Abidjan"      "Africa/Accra"        "Africa/Addis_Ababa"  
## [4] "Africa/Algiers"     "Africa/Asmara"       "Africa/Asmera"
```

```
(x1 <- ymd_hms("2015-06-01 12:00:00", tz = "America/New_York"))  
(x2 <- ymd_hms("2015-06-01 18:00:00", tz = "Europe/Copenhagen"))  
(x3 <- ymd_hms("2015-06-02 04:00:00", tz = "Pacific/Auckland"))
```

```
## [1] "2015-06-01 12:00:00 EDT"  
## [1] "2015-06-01 18:00:00 CEST"  
## [1] "2015-06-02 04:00:00 NZST"
```

```
x1 - x2  
x1 - x3
```

```
## Time difference of 0 secs  
## Time difference of 0 secs
```

UTC:

```
x4 <- c(x1, x2, x3)  
x4
```

```
## [1] "2015-06-01 18:00:00 CEST" "2015-06-01 18:00:00 CEST"  
## [3] "2015-06-01 18:00:00 CEST"
```


Keep the instant in time:

```
x4a <- with_tz(x4, tzone = "Australia/Lord_Howe")  
x4a  
x4a - x4
```

```
## [1] "2015-06-02 02:30:00 +1030" "2015-06-02 02:30:00 +1030"  
## [3] "2015-06-02 02:30:00 +1030"  
## Time differences in secs  
## [1] 0 0 0
```

Change the instant in time:

```
x4b <- force_tz(x4, tzone = "Australia/Lord_Howe")  
x4b  
x4b - x4
```

```
## [1] "2015-06-01 16:00:00 +1030" "2015-06-01 16:00:00 +1030"  
## [3] "2015-06-01 16:00:00 +1030"  
## Time differences in hours  
## [1] -10.5 -10.5 -10.5
```

1 Relational data

2 Combining tables

3 Dates and times

4 Factors

5 Strings

- Used to work with categorical variables (i.e., that have a fixed and known set of possible values.
- Useful to display character vectors in a non-alphabetical order.

The **forcats** package:

- Range of helpers for working with factors.
- Not part of the core tidyverse, so we need to load it explicitly.

```
library(forcats)
```

Imagine that you have a variable that records month:

```
x1 <- c("Dec", "Apr", "Jan", "Mar")
```

Using a string to record this variable has two problems:

1. Twelve possible months and nothing saving you from typos.
2. It doesn't sort in a useful way:

```
x2 <- c("Dec", "Apr", "Jam", "Mar")  
sort(x1)
```

```
## [1] "Apr" "Dec" "Jan" "Mar"
```

Start by creating a list of the valid **levels**:

```
month_levels <- c("Jan", "Feb", "Mar", "Apr", "May", "Jun",  
                  "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")
```

Then create a factor:

```
y1 <- factor(x1, levels = month_levels)  
y1
```

```
## [1] Dec Apr Jan Mar  
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

```
sort(y1)
```

```
## [1] Jan Mar Apr Dec  
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

```
factor(x1) ## without levels
```

```
## [1] Dec Apr Jan Mar  
## Levels: Apr Dec Jan Mar
```

Notice:

```
y2 <- factor(x2, levels = month_levels)
y2
```

```
## [1] Dec Apr <NA> Mar
## Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

Other ordering:

```
f1 <- factor(x1, levels = unique(x1))
f1

f2 <- x1 %>% factor() %>% fct_inorder()
f2
```

```
## [1] Dec Apr Jan Mar
## Levels: Dec Apr Jan Mar
## [1] Dec Apr Jan Mar
## Levels: Dec Apr Jan Mar
```

To access the set of valid levels directly: `levels(f2)`.

Sample from the General Social Survey:

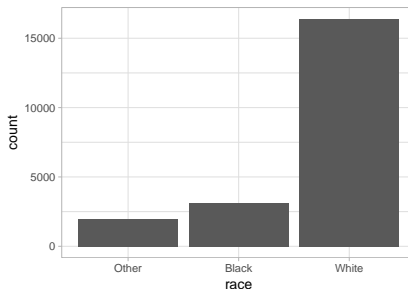
```
gss_cat
```

```
## # A tibble: 21,483 x 9
##   year marital    age race  rincome partyid  relig  denom  tvhours
##   <int> <fct>    <int> <fct> <fct>    <fct>  <fct> <fct>    <int>
## 1  2000 Never m~    26 White $8000 t~ Ind,nea~ Prote~ South~    12
## 2  2000 Divorced    48 White $8000 t~ Not str~ Prote~ Bapti~    NA
## 3  2000 Widowed    67 White Not app~ Indepen~ Prote~ No de~     2
## 4  2000 Never m~    39 White Not app~ Ind,nea~ Ortho~ Not a~     4
## 5  2000 Divorced    25 White Not app~ Not str~ None   Not a~     1
## 6  2000 Married    25 White $20000 ~ Strong ~ Prote~ South~    NA
## 7  2000 Never m~    36 White $25000 ~ Not str~ Chris~ Not a~     3
## 8  2000 Divorced    44 White $7000 t~ Ind,nea~ Prote~ Luthe~    NA
## 9  2000 Married    44 White $25000 ~ Not str~ Prote~ Other     0
## 10 2000 Married    47 White $25000 ~ Strong ~ Prote~ South~     3
## # ... with 21,473 more rows
```

More info with `?gss_cat`.

Levels of a factor stored in a tibble

```
ggplot(gss_cat, aes(race)) + geom_bar()
```

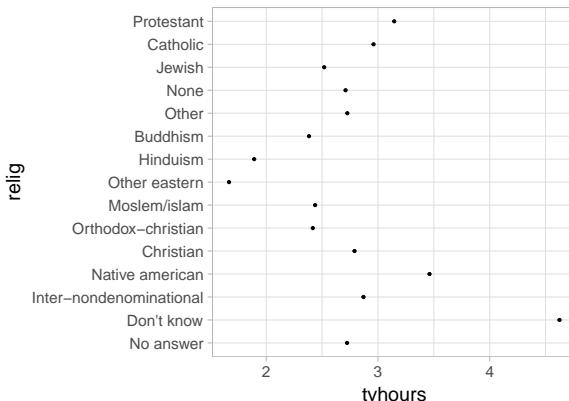


```
gss_cat %>% count(race)
```

```
## # A tibble: 3 x 2
##   race      n
##   <fct> <int>
## 1 Other  1959
## 2 Black  3129
## 3 White 16395
```

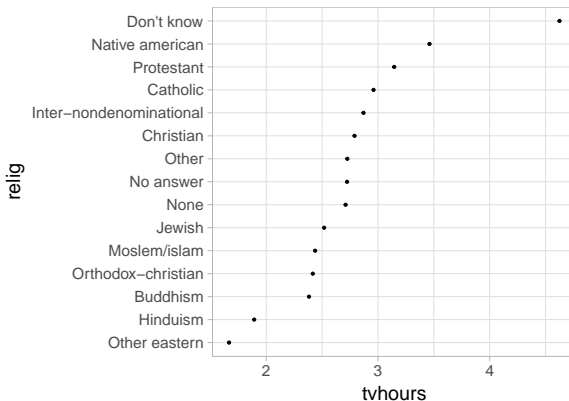

What's wrong here?

```
relig_summary <- gss_cat %>%  
  group_by(relig) %>%  
  summarise(age = mean(age, na.rm = TRUE),  
            tvhours = mean(tvhours, na.rm = TRUE),  
            n = n())  
ggplot(relig_summary, aes(tvhours, relig)) + geom_point()
```



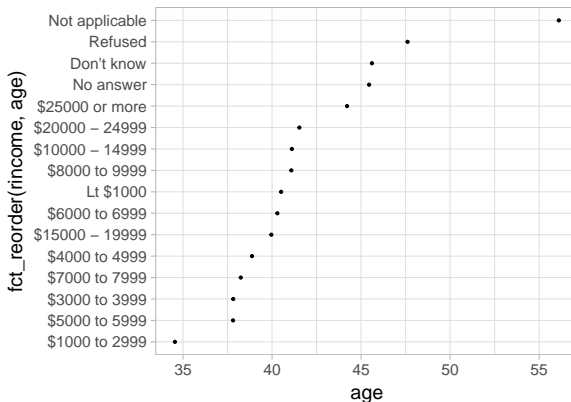
Modifying factor order

```
relig_summary %>%  
  mutate(relig = fct_reorder(relig, tvhours)) %>%  
  ggplot(aes(tvhours, relig)) + geom_point()
```



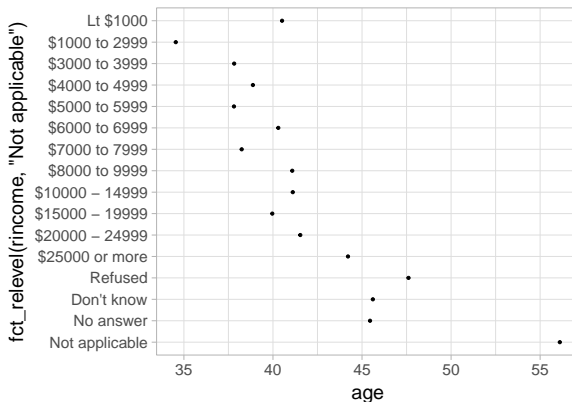
What's wrong here?

```
rincome_summary <- gss_cat %>%  
  group_by(rincome) %>%  
  summarise(age = mean(age, na.rm = TRUE),  
            tvhours = mean(tvhours, na.rm = TRUE),  
            n = n())  
ggplot(rincome_summary, aes(age, fct_reorder(rincome, age))) + geom_point()
```



Modify factor order II

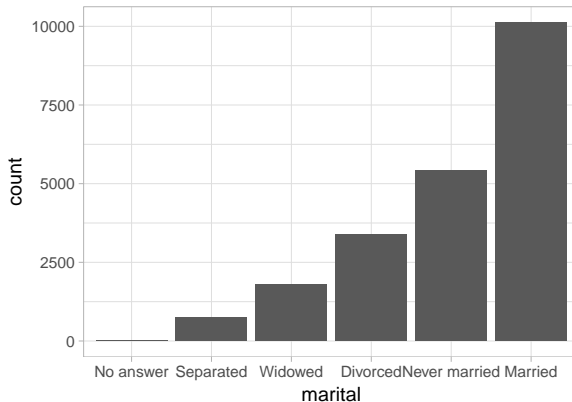
```
ggplot(rincome_summary, aes(age, fct_relevel(rincome,  
                                             "Not applicable")))) +  
  geom_point()
```



Why do you think the average age for “Not applicable” is so high?

Modify factor order III

```
gss_cat %>%  
  mutate(marital = marital %>% fct_infreq() %>% fct_rev()) %>%  
  ggplot(aes(marital)) + geom_bar()
```



More powerful than changing the orders of the levels is changing their values:

- To clarify labels for publication.
- To collapse levels for high-level displays.

What's wrong here?

```
gss_cat %>% count(partyid)
```

```
## # A tibble: 10 x 2
##   partyid          n
##   <fct>         <int>
## 1 No answer      154
## 2 Don't know      1
## 3 Other party    393
## 4 Strong republican 2314
## 5 Not str republican 3032
## 6 Ind,near rep    1791
## 7 Independent    4119
## 8 Ind,near dem    2499
## 9 Not str democrat 3690
## 10 Strong democrat 3490
```

Modifying factor levels II

```
gss_cat %>%  
  mutate(partyid = fct_recode(partyid,  
    "Republican, strong"      = "Strong republican",  
    "Republican, weak"       = "Not str republican",  
    "Independent, near rep"  = "Ind,near rep",  
    "Independent, near dem"  = "Ind,near dem",  
    "Democrat, weak"        = "Not str democrat",  
    "Democrat, strong"      = "Strong democrat")) %>%  
  count(partyid)
```

```
## # A tibble: 10 x 2  
##   partyid      n  
##   <fct>      <int>  
## 1 No answer    154  
## 2 Don't know     1  
## 3 Other party   393  
## 4 Republican, strong 2314  
## 5 Republican, weak  3032  
## 6 Independent, near rep 1791  
## 7 Independent    4119  
## 8 Independent, near dem 2499  
## 9 Democrat, weak   3690  
## 10 Democrat, strong 3490
```



```
gss_cat %>%  
  mutate(partyid = fct_recode(partyid,  
    "Republican, strong"      = "Strong republican",  
    "Republican, weak"       = "Not str republican",  
    "Independent, near rep"  = "Ind,near rep",  
    "Independent, near dem"  = "Ind,near dem",  
    "Democrat, weak"        = "Not str democrat",  
    "Democrat, strong"      = "Strong democrat",  
    "Other"                 = "No answer",  
    "Other"                 = "Don't know",  
    "Other"                 = "Other party" )) %>% count(partyid)
```

```
## # A tibble: 8 x 2  
##   partyid      n  
##   <fct>      <int>  
## 1 Other      548  
## 2 Republican, strong 2314  
## 3 Republican, weak  3032  
## 4 Independent, near rep 1791  
## 5 Independent      4119  
## 6 Independent, near dem 2499  
## 7 Democrat, weak    3690  
## 8 Democrat, strong  3490
```

```
gss_cat %>%  
  mutate(partyid = fct_collapse(partyid,  
    other = c("No answer", "Don't know", "Other party"),  
    rep = c("Strong republican", "Not str republican"),  
    ind = c("Ind,near rep", "Independent", "Ind,near dem"),  
    dem = c("Not str democrat", "Strong democrat")  
  )) %>%  
  count(partyid)
```

```
## # A tibble: 4 x 2  
##   partyid      n  
##   <fct>    <int>  
## 1 other      548  
## 2 rep      5346  
## 3 ind      8409  
## 4 dem      7180
```

Collapsing factor III

```
gss_cat %>%  
  mutate(relig = fct_lump(relig)) %>%  
  count(relig)
```

```
## # A tibble: 2 x 2  
##   relig      n  
##   <fct>    <int>  
## 1 Protestant 10846  
## 2 Other      10637
```

```
gss_cat %>%  
  mutate(relig = fct_lump(relig, n = 3)) %>%  
  count(relig, sort = TRUE)
```

```
## # A tibble: 4 x 2  
##   relig      n  
##   <fct>    <int>  
## 1 Protestant 10846  
## 2 Catholic    5124  
## 3 None       3523  
## 4 Other      1990
```

1 Relational data

2 Combining tables

3 Dates and times

4 Factors

5 Strings

```
library(stringr) # package for string manipulation

# To create strings
string1 <- "This is a string"
string2 <- 'To get a "quote" inside a string, use single quotes'
```

Backslash as escape character:

```
double_quote <- "\"\" # or '\"'
single_quote <- '\'\'' # or '\"'
```

The printed representation is not the string itself:

```
x <- c("\"", "\\")
x
writeLines(x)
```

```
## [1] "\" " \" \"
## "
## \
```

Special characters:

- Use `"\n"`, for newline, or `"\t"`, for tab.
- Complete list by requesting help on `"` (`?'"'`, or `?'"'`)

Other useful things:

```
(x <- "\u00b5") # Non-English characters
```








```
## [1] "µ"
```

```
c("one", "two", "three") # Character vectors
```

```
## [1] "one"    "two"    "three"
```

```
str_length(c("a", "R for data science", NA)) # String length
```

```
## [1]  1 18 NA
```

```
>  str_c      {stringr}  
>  str_conv   {stringr}  
>  str_count  {stringr}  
>  str_detect {stringr}  
>  str_dup    {stringr}  
>  str_extract {stringr}  
>  str_extract_all {stringr}
```

```
> str_|
```

```
str_c(..., sep = "", collapse = NULL)
```

To understand how `str_c` works, you need to imagine that you are building up a matrix of strings. Each input argument forms a column, and is expanded to the length of the longest argument, using the usual recycling rules. The `sep` string is inserted between each column. If `collapse` is `NULL` each row is collapsed into a single string. If non-`NULL` that string is inserted at the end of each row, and the entire matrix collapsed to a single string.

Press F1 for additional help

Combining strings:

```
str_c("x", "y")  
str_c("x", "y", "z")  
str_c("x", "y", sep = ", ")
```

```
## [1] "xy"  
## [1] "xyz"  
## [1] "x, y"
```

Missing values:

```
x <- c("abc", NA)  
str_c("|-", x, "-|")
```

```
## [1] "|-abc-|" NA
```

```
str_c("|-", str_replace_na(x), "-|")
```

```
## [1] "|-abc-|" "|-NA-|"
```


Recycling:

```
str_c("prefix-", c("a", "b", "c"), "-suffix")
```

```
## [1] "prefix-a-suffix" "prefix-b-suffix" "prefix-c-suffix"
```

Collapsing a vector of strings:

```
str_c(c("x", "y", "z"), collapse = ", ")
```

```
## [1] "x, y, z"
```

Subsetting strings

```
x <- c("Apple", "Banana", "Pear")  
str_sub(x, 1, 3)
```

```
## [1] "App" "Ban" "Pea"
```

```
str_sub(x, -3, -1)
```

```
## [1] "ple" "ana" "ear"
```

```
str_sub("a", 1, 5)
```

```
## [1] "a"
```

```
str_sub(x, 1, 1) <- str_to_lower(str_sub(x, 1, 1))  
x
```

```
## [1] "apple" "banana" "pear"
```

See also `str_to_upper()` or `str_to_title()`.

```
# Turkish has two i's: with and without a dot, and it  
# has a different rule for capitalising them:  
str_to_upper(c("i", "I"))
```

```
## [1] "I" "I"
```

```
str_to_upper(c("i", "I"), locale = "tr")
```

```
## [1] "İ" "I"
```

The locale:

- An ISO 639 language code, which is a **two or three letter abbreviation**
- If blank, R uses the current locale, as provided by your operating system.

A language that allows you to describe patterns in strings.

Allows you for instance to:

- Determine which strings match a pattern.
- Find the positions of matches.
- Extract the content of matches.
- Replace matches with new values.
- Split a string based on a match.

[Read the chapter on strings from the book!](#)