

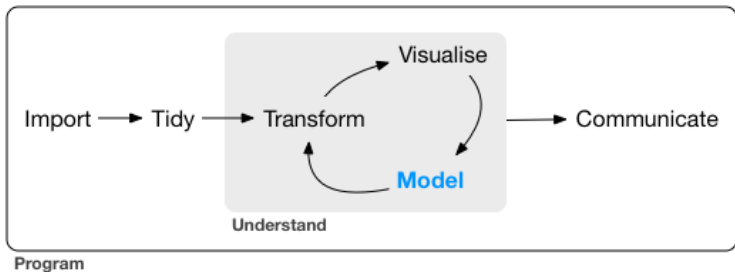
Lecture 7: Modeling I

Data Science for Business Analytics

Thibault Vatter

Department of Statistics, Columbia University and HEC Lausanne, UNIL

4/9/2018



- This morning:
 - ▶ how models work mechanistically (focus on linear models),
 - ▶ how to use models to find patterns in real data.
- This afternoon:
 - ▶ how to use **many** simple models,
 - ▶ how to combine modeling and programming tools.

As usual, most of the material is borrowed from [R for data science](#).

1 Model basics

2 Model building

1. Each observation can either be used for exploration **OR** confirmation, not both.
2. You can use an observation
 - ▶ as many times as you like for exploration,
 - ▶ only once for confirmation.

When using an observation twice, switch from confirmation to exploration.

Goals:

- Provide a simple low-dimensional summary of a dataset.
- Often partition data into patterns and residuals.
- Help peel back layers of structure (since strong patterns hide subtler trends).

Two parts to a model:

1. **Family of models:** a precise, but generic, pattern to capture.
 - ▶ A straight line, or a quadratic curve.
 - ▶ Equations like $y = a_1 * x + a_2$ or $y = a_1 * x^a_2$ (with x and y known variables and a_1 and a_2 parameters).
2. **Fitted model:** member of the family that is closest to the data.
 - ▶ $y = 3 * x + 7$ or $y = 9 * x^2$.

All models are wrong, but some are useful. —George Box


- A fitted model is “just” the closest model to the data from a family of models.
- The “best” model (according to some criteria):
 - ▶ isn't necessarily a good model,
 - ▶ isn't necessarily “true”.

The goal is not to uncover truth, but to discover useful approximations.

```
library(tidyverse)
library(modelr)
```


A simulated dataset


```
ggplot(sim1, aes(x, y)) + geom_point(size = 2)
```



Linear family?

```
models <- tibble(a1 = runif(250, -20, 40),  
                 a2 = runif(250, -5, 5))  
ggplot(sim1, aes(x, y)) + geom_point(size = 2) +  
  geom_abline(aes(intercept = a1, slope = a2),  
             data = models, alpha = 1/4)
```





This distance is the difference between

- the y value given by the model (the **prediction**),
- and the actual y value in the data (the **response**).

The model family:

```
model1 <- function(a, data) a[1] + data$x * a[2]
```

```
model1(c(7, 1.5), sim1)
```

```
## [1] 8.5 8.5 8.5 10.0 10.0 10.0 11.5 11.5 11.5 13.0 13.0 13.0 14.5  
## [14] 14.5 14.5 16.0 16.0 16.0 17.5 17.5 17.5 19.0 19.0 19.0 20.5 20.5  
## [27] 20.5 22.0 22.0 22.0
```

Root-mean-square error (RMSE):

```
measure_distance <- function(mod, data) {  
  diff <- data$y - model1(mod, data)  
  sqrt(mean(diff ^ 2))  
}
```

```
measure_distance(c(7, 1.5), sim1)
```

```
## [1] 2.665212
```

RMSE for each model


```
sim1_dist <- function(a1, a2) measure_distance(c(a1, a2), sim1)

(models <- models %>% mutate(dist = map2_dbl(a1, a2, sim1_dist)))
```

```
## # A tibble: 250 x 3
##       a1      a2  dist
##   <dbl>  <dbl> <dbl>
## 1  33.8    4.43  43.3
## 2  -4.07    2.62   5.82
## 3   2.33    4.33  12.7
## 4  14.4   -0.293   7.58
## 5  34.5    1.04  24.9
## 6  -7.90  -0.150  25.1
## 7  33.9   -3.91  17.5
## 8  36.7   -2.52  15.2
## 9  19.6   -0.0149  7.50
## 10 17.7   -1.27  10.9
## # ... with 240 more rows
```


The 10 best models

```
ggplot(sim1, aes(x, y)) +  
  geom_abline(aes(intercept = a1, slope = a2, color = -dist),  
             data = filter(models, rank(dist) <= 10)) +  
  geom_point(size = 2)
```




Grid search

```
grid <- expand.grid(a1 = seq(-5, 20, length = 25),  
                  a2 = seq(1, 3, length = 25)) %>%  
  mutate(dist = map2_dbl(a1, a2, sim1_dist))  
  
ggplot(grid, aes(a1, a2)) + geom_point(aes(color = -dist)) +  
  geom_point(data = filter(grid, rank(dist) <= 10),  
            size = 4, color = "red")
```



Grid search cont'd

```
ggplot(sim1, aes(x, y)) +  
  geom_abline(aes(intercept = a1, slope = a2, color = -dist),  
             data = filter(grid, rank(dist) <= 10)) +  
  geom_point(size = 2)
```




Newton-Raphson search

```
(best <- optim(c(0, 0), measure_distance, data = sim1)$par)
```

```
## [1] 4.222248 2.051204
```

```
ggplot(sim1, aes(x, y)) + geom_point(size = 2) +  
  geom_abline(intercept = best[1], slope = best[2])
```




- General form is e.g. $y = a_1 + a_2 * x_1 + a_3 * x_2 + \dots + a_n * x_{(n-1)}$
- `lm()` specify the model family using formulas (e.g., $y \sim x$, which `lm()` translates to a function like $y = a_1 + a_2 * x$)

```
sim1_mod <- lm(y ~ x, data = sim1)
coef(sim1_mod)
```

```
## (Intercept)          x
##    4.220822    2.051533
```


What is going on here ?

```
sim1a <- tibble(x = rep(1:10, each = 3),  
               y = 6 + x * 1.5 + rnorm(length(x)))  
sim1a$y[sim1a$x == max(sim1a$x)] <- 40  
  
sim1a_mod <- lm(y ~ x, data = sim1a)  
ggplot(sim1a, aes(x, y)) + geom_point(size = 2) +  
  geom_abline(intercept = coef(sim1a_mod)[1], slope = coef(sim1a_mod)[2])
```



Two interesting quantities to look at:

- predictions,
- and residuals.

Visualizing predictions: step 1/3

```
(grid <- sim1 %>% modelr::data_grid(x))
```

```
## # A tibble: 10 x 1
```

```
##       x
```

```
##   <int>
```

```
## 1     1
```

```
## 2     2
```

```
## 3     3
```

```
## 4     4
```

```
## 5     5
```

```
## 6     6
```

```
## 7     7
```

```
## 8     8
```

```
## 9     9
```

```
## 10    10
```


Visualizing predictions: step 2/3

```
(grid <- grid %>% add_predictions(sim1_mod))
```

```
## # A tibble: 10 x 2
##       x   pred
##   <int> <dbl>
## 1     1    6.27
## 2     2    8.32
## 3     3   10.4
## 4     4   12.4
## 5     5   14.5
## 6     6   16.5
## 7     7   18.6
## 8     8   20.6
## 9     9   22.7
## 10    10   24.7
```

Visualizing predictions: step 3/3

```
ggplot(sim1, aes(x, y)) + geom_point(size = 2) +  
  geom_line(aes(y = pred), data = grid, size = 1)
```





Visualizing residuals: step 1/2

```
(sim1 <- sim1 %>% add_residuals(sim1_mod))
```

```
## # A tibble: 30 x 3
##       x     y   resid
##   <int> <dbl>   <dbl>
## 1     1     1  4.20 -2.07
## 2     1     1  7.51  1.24
## 3     1     1  2.13 -4.15
## 4     2     2  8.99  0.665
## 5     2    10.2  1.92
## 6     2    11.3  2.97
## 7     3     3  7.36 -3.02
## 8     3    10.5  0.130
## 9     3    10.5  0.136
## 10    4    12.4  0.00763
## # ... with 20 more rows
```

Visualizing residuals: step 2/2

```
ggplot(sim1, aes(resid)) + geom_freqpoly(binwidth = 0.5)  
  
ggplot(sim1, aes(x, resid)) + geom_ref_line(h = 0) + geom_point()
```



1. Instead of using `lm()` to fit a straight line, you can use `loess()` to fit a smooth curve. Repeat the process of model fitting, grid generation, predictions, and visualisation on `sim1` using `loess()` instead of `lm()`. How does the result compare to `geom_smooth()`?
2. `add_predictions()` is paired with `gather_predictions()` and `spread_predictions()`. How do these three functions differ?
3. What does `geom_ref_line()` do? What package does it come from? Why is displaying a reference line in plots showing residuals useful and important?
4. Why might you want to look at a frequency polygon of absolute residuals? What are the pros and cons compared to looking at the raw residuals?

- A way of getting “special behavior”.
- “Capture variables” so they can be interpreted by the function.
- Sometimes called “Wilkinson-Rogers notation” from [Symbolic Description of Factorial Models for Analysis of Variance](#)

Behind the scenes:

```
df <- tribble(~y, ~x1, ~x2,  
              4, 2, 5,  
              5, 1, 6)  
model_matrix(df, y ~ x1)
```

```
## # A tibble: 2 x 2  
##   `(Intercept)`    x1  
##           <dbl> <dbl>  
## 1           1.00  2.00  
## 2           1.00  1.00
```

Without intercept:

```
model_matrix(df, y ~ x1 - 1)
```

```
## # A tibble: 2 x 1
##       x1
##   <dbl>
## 1  2.00
## 2  1.00
```

Adding a second variable:

```
model_matrix(df, y ~ x1 + x2)
```

```
## # A tibble: 2 x 3
##   `(Intercept)`    x1    x2
##   <dbl> <dbl> <dbl>
## 1      1.00  2.00  5.00
## 2      1.00  1.00  6.00
```


```
df <- tribble(~ sex, ~ response,  
              "male", 1,  
              "female", 2,  
              "male", 1)  
model_matrix(df, response ~ sex)
```

```
## # A tibble: 3 x 2  
##   `(Intercept)` sexmale  
##           <dbl>   <dbl>  
## 1           1.00     1.00  
## 2           1.00      0  
## 3           1.00     1.00
```

Why doesn't R also create a sexfemale column?

Another simulated dataset

```
ggplot(sim2, aes(x, y)) + geom_point(size = 2)
```




```
mod2 <- lm(y ~ x, data = sim2)

(grid <- sim2 %>% data_grid(x) %>% add_predictions(mod2))

## # A tibble: 4 x 2
##   x      pred
##   <chr> <dbl>
## 1 a      1.15
## 2 b      8.12
## 3 c      6.13
## 4 d      1.91
```

Visualize the results

```
ggplot(sim2, aes(x)) + geom_point(aes(y = y), size = 2) +  
  geom_point(data = grid, aes(y = pred), color = "red", size = 4)
```




What's happening here?

```
tibble(x = "e") %>% add_predictions(mod2)
```

```
## Error in model.frame.default(Terms, newdata, na.action = na.action, xlev = o
```

Interactions (cont. and cat.)

```
ggplot(sim3, aes(x1, y)) + geom_point(aes(color = x2))
```



Two possible models

```
mod1 <- lm(y ~ x1 + x2, data = sim3)
mod2 <- lm(y ~ x1 * x2, data = sim3)
```

Note that:

- $y \sim x1 + x2$ is translated to $y = a_0 + a_1 * x1 + a_2 * x2$.
- $y \sim x1 * x2$ is translated to $y = a_0 + a_1 * x1 + a_2 * x2 + a_{12} * x1 * x2$.

Two new tricks to visualize them


1. Give `data_grid()` both variables.
2. To generate predictions from both models simultaneously, use
 - ▶ `gather_predictions()` to add predictions as rows,
 - ▶ or `spread_predictions()` to add predictions as columns.

```
(grid <- sim3 %>% data_grid(x1, x2) %>% gather_predictions(mod1, mod2))
```

```
## # A tibble: 80 x 4
##   model    x1 x2    pred
##   <chr> <int> <fct> <dbl>
## 1 mod1     1 a     1.67
## 2 mod1     1 b     4.56
## 3 mod1     1 c     6.48
## 4 mod1     1 d     4.03
## 5 mod1     2 a     1.48
## 6 mod1     2 b     4.37
## 7 mod1     2 c     6.28
## 8 mod1     2 d     3.84
## 9 mod1     3 a     1.28
## 10 mod1    3 b     4.17
## # ... with 70 more rows
```

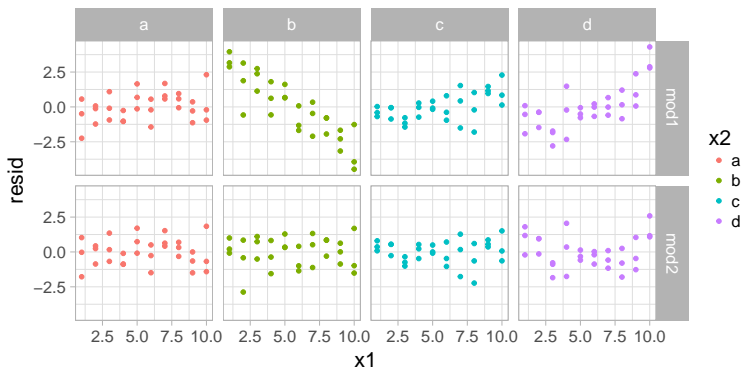
Using facetting

```
ggplot(sim3, aes(x1, y, color = x2)) + geom_point() +  
  geom_line(data = grid, aes(y = pred)) + facet_wrap(~ model)
```



Which model is better?

```
sim3 <- sim3 %>% gather_residuals(mod1, mod2)  
  
ggplot(sim3, aes(x1, resid, color = x2)) + geom_point() +  
  facet_grid(model ~ x2)
```



Interactions (two continuous)


```
mod1 <- lm(y ~ x1 + x2, data = sim4)
mod2 <- lm(y ~ x1 * x2, data = sim4)

(grid <- sim4 %>%
  data_grid(x1 = seq_range(x1, 5), x2 = seq_range(x2, 5)) %>%
  gather_predictions(mod1, mod2))
```

```
## # A tibble: 50 x 4
##   model    x1    x2  pred
##   <chr> <dbl> <dbl> <dbl>
## 1 mod1 -1.00 -1.00  0.996
## 2 mod1 -1.00 -0.500 -0.395
## 3 mod1 -1.00  0      -1.79
## 4 mod1 -1.00  0.500 -3.18
## 5 mod1 -1.00  1.00  -4.57
## 6 mod1 -0.500 -1.00  1.91
## 7 mod1 -0.500 -0.500  0.516
## 8 mod1 -0.500  0      -0.875
## 9 mod1 -0.500  0.500 -2.27
## 10 mod1 -0.500  1.00  -3.66
## # ... with 40 more rows
```


?seq_range for other arguments (e.g., pretty = TRUE for tables).

```
ggplot(grid, aes(x1, x2)) + geom_tile(aes(fill = pred)) +  
  facet_wrap(~ model)
```




Slices with respect to x2

```
ggplot(grid, aes(x1, pred, color = x2, group = x2)) + geom_line() +  
  facet_wrap(~ model)
```



Slices with respect to x1

```
ggplot(grid, aes(x2, pred, color = x1, group = x1)) + geom_line() +  
  facet_wrap(~ model)
```



- $\log(y) \sim \sqrt{x_1} + x_2$ is transformed to $\log(y) = a_1 + a_2 * \sqrt{x_1} + a_3 * x_2$.
- If the transformation involves $+$, $*$, $^$, or $-$, wrap it in $I()$:
 - ▶ $y \sim x + I(x^2) \equiv y = a_1 + a_2 * x + a_3 * x^2$.
 - ▶ $y \sim x^2 + x \equiv y \sim x * x + x \equiv y = a_1 + a_2 * x$.

```
df <- tribble(~y, ~x, 1, 1, 2, 2, 3, 3)
model_matrix(df, y ~ x^2 + x)
model_matrix(df, y ~ I(x^2) + x)
```

```
## # A tibble: 3 x 2
##   `(Intercept)`      x
##   <dbl> <dbl>
## 1      1.00  1.00
## 2      1.00  2.00
## 3      1.00  3.00
## # A tibble: 3 x 3
##   `(Intercept)` `I(x^2)`      x
##   <dbl> <dbl> <dbl>
## 1      1.00      1.00  1.00
## 2      1.00      4.00  2.00
## 3      1.00      9.00  3.00
```

To get $y = a_1 + a_2 * x + a_3 * x^2$:

```
model_matrix(df, y ~ poly(x, 2))
```

```
## # A tibble: 3 x 3
##   `(Intercept)`      `poly(x, 2)1` `poly(x, 2)2`
##   <dbl>            <dbl>      <dbl>
## 1      1.00 -0.707              0.408
## 2      1.00 -0.000000000000000785 -0.816
## 3      1.00  0.707              0.408
```


```
library(splines)
model_matrix(df, y ~ ns(x, 2))
```

```
## # A tibble: 3 x 3
##   `(Intercept)` `ns(x, 2)1` `ns(x, 2)2`
##   <dbl>         <dbl>         <dbl>
## 1         1.00         0             0
## 2         1.00        0.566        -0.211
## 3         1.00        0.344         0.771
```

A non-linear function

```
sim5 <- tibble(x = seq(0, 3.5 * pi, length = 50),  
              y = 4 * sin(x) + rnorm(length(x)))
```

```
ggplot(sim5, aes(x, y)) + geom_point()
```




Three models using splines:

```
mod1 <- lm(y ~ ns(x, 1), data = sim5)  
mod2 <- lm(y ~ ns(x, 2), data = sim5)  
mod5 <- lm(y ~ ns(x, 5), data = sim5)
```

```
grid <- sim5 %>%
  data_grid(x = seq_range(x, n = 50, expand = 0.1)) %>%
  gather_predictions(mod1, mod2, mod5, .pred = "y")

ggplot(sim5, aes(x, y)) + geom_point() +
  geom_line(data = grid, color = "red") +
  facet_wrap(~ model)
```



```
options(na.action = na.warn)
df <- tribble(~x, ~y,
              1, 2.2,
              2, NA,
              3, 3.5,
              4, 8.3,
              NA, 10)
```

```
mod <- lm(y ~ x, data = df)
```

```
## Warning: Dropping 2 rows with missing values
```

```
mod <- lm(y ~ x, data = df, na.action = na.exclude)
nobs(mod)
```

```
## [1] 3
```

- **Generalised linear models**, e.g. `stats::glm()`:
 - ▶ LMs assume continuous responses and Gaussian errors.
 - ▶ GLMs extend LMs to other distributions, including non-continuous responses (e.g. binary data or counts).
- **Generalised additive models**, e.g. `mgcv::gam()`:
 - ▶ Extend GLMs to incorporate arbitrary smooth functions.
 - ▶ A formula like $y \sim s(x)$ becomes an equation like $y = f(x)$.
- **Penalized linear models**, e.g. `glmnet::glmnet()`:
 - ▶ Add penalties to favor simpler models.
 - ▶ “Generalize” better to new datasets.
- **Robust linear models**, e.g. `MASS::rlm()`:
 - ▶ Tweaks distance to downweight outliers.
 - ▶ Less sensitive to outliers, but slightly worse without outliers.
- **Trees**, e.g. `rpart::rpart()`:
 - ▶ A piece-wise constant model splitting the data into small pieces.
 - ▶ Powerful when aggregated as **random forests** (e.g. `randomForest::randomForest()`) or **gradient boosting machines** (e.g. `xgboost::xgboost()`).

1 Model basics

2 Model building

To partition data into pattern and residuals:


1. Find patterns with visualisation.
2. Make them concrete and precise with a model.
3. Repeat 1. and 2. after replacing the old response variable with the residuals from the model.

How about large and complex datasets?

- ML approaches “simply” focus on predictive ability.
- Issues:
 - ▶ black boxes,
 - ▶ (sometimes) hard to use domain knowledge,
 - ▶ (often) difficult to assess whether or not the model will continue to work in the long-term
- Usually, combinations of both approaches are preferred.

The diamonds dataset


```
ggplot(diamonds, aes(cut, price)) + geom_boxplot()  
ggplot(diamonds, aes(color, price)) + geom_boxplot()  
ggplot(diamonds, aes(clarity, price)) + geom_boxplot()
```



Why are low quality diamonds more expensive?

Price and carat

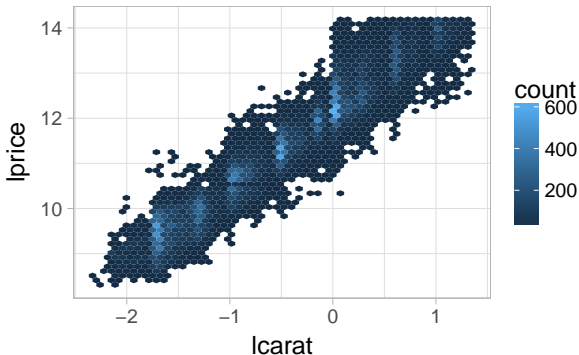
```
ggplot(diamonds, aes(carat, price)) + geom_hex(bins = 50)
```



A couple of tweaks


1. Focus on diamonds < 2.5 carats (99.7% of the data).
2. Log-transform the carat and price.

```
diamonds2 <- diamonds %>% filter(carat <= 2.5) %>%  
  mutate(lprice = log2(price), lcarat = log2(carat))  
  
ggplot(diamonds2, aes(lcarat, lprice)) + geom_hex(bins = 50)  
mod_diamond <- lm(lprice ~ lcarat, data = diamonds2)
```




Visualize the predictions

```
grid <- diamonds2 %>% data_grid(carat = seq_range(carat, 20)) %>%  
  mutate(lcarat = log2(carat)) %>%  
  add_predictions(mod_diamond, "lprice") %>%  
  mutate(price = 2 ^ lprice)  
  
ggplot(diamonds2, aes(carat, price)) + geom_hex(bins = 50) +  
  geom_line(data = grid, color = "red", size = 1)
```




Visualize the residuals

```
diamonds2 <- diamonds2 %>% add_residuals(mod_diamond, "lresid")  
  
ggplot(diamonds2, aes(lcarat, lresid)) + geom_hex(bins = 50)
```



Replace price by residuals

```
ggplot(diamonds2, aes(cut, lresid)) + geom_boxplot()  
ggplot(diamonds2, aes(color, lresid)) + geom_boxplot()  
ggplot(diamonds2, aes(clarity, lresid)) + geom_boxplot()
```



A more complicated model


```
mod_diamond2 <- lm(lprice ~ lcarat + color + cut + clarity,  
                  data = diamonds2)
```

```
(grid <- diamonds2 %>%  
  data_grid(cut, lcarat = -0.515,  
            color = "G", clarity = "SI1") %>%  
  add_predictions(mod_diamond2))
```

```
## # A tibble: 5 x 5  
##   cut      lcarat color clarity  pred  
##   <ord>    <dbl> <chr>  <chr>  <dbl>  
## 1 Fair      -0.515 G      SI1     11.0  
## 2 Good      -0.515 G      SI1     11.1  
## 3 Very Good -0.515 G      SI1     11.2  
## 4 Premium   -0.515 G      SI1     11.2  
## 5 Ideal     -0.515 G      SI1     11.2
```


Visualize the residuals

```
diamonds2 <- diamonds2 %>%  
  add_residuals(mod_diamond2, "lresid2")  
  
ggplot(diamonds2, aes(lcarat, lresid2)) +  
  geom_hex(bins = 50)
```




The number of daily flights


```
daily <- flights %>%  
  mutate(date = make_date(year, month, day)) %>%  
  group_by(date) %>%  
  summarise(n = n())
```

What affects this number?

```
ggplot(daily, aes(date, n)) + geom_line()
```




```
daily <- daily %>% mutate(wday = wday(date, label = TRUE))  
  
ggplot(daily, aes(wday, n)) + geom_boxplot()  
  
mod <- lm(n ~ wday, data = daily)
```



Visualize the predictions


```
grid <- daily %>% data_grid(wday) %>%  
  add_predictions(mod, "n")  
  
ggplot(daily, aes(wday, n)) + geom_boxplot() +  
  geom_point(data = grid, color = "red", size = 4)
```



Visualize the residuals


```
daily <- daily %>% add_residuals(mod)

daily %>% ggplot(aes(date, resid)) +
  geom_ref_line(h = 0) + geom_line()
```



What happens here?

```
ggplot(daily, aes(date, resid, colour = wday)) +  
  geom_ref_line(h = 0) + geom_line()
```




What happens here?

```
daily %>% filter(resid < -100)
```

```
## # A tibble: 11 x 4
##   date           n wday  resid
##   <date>       <int> <ord> <dbl>
## 1 2013-01-01    842 Tue   -109
## 2 2013-01-20    786 Sun   -105
## 3 2013-05-26    729 Sun   -162
## 4 2013-07-04    737 Thu   -229
## 5 2013-07-05    822 Fri   -145
## 6 2013-09-01    718 Sun   -173
## 7 2013-11-28    634 Thu   -332
## 8 2013-11-29    661 Fri   -306
## 9 2013-12-24    761 Tue   -190
##10 2013-12-25    719 Wed   -244
##11 2013-12-31    776 Tue   -175
```



What happens here?

```
daily %>%  
  ggplot(aes(date, resid)) +  
  geom_ref_line(h = 0) +  
  geom_line(colour = "grey50") +  
  geom_smooth(se = FALSE, span = 0.20)
```



Seasonal Saturday effect


```
daily %>% filter(wday == "Sat") %>%  
  ggplot(aes(date, n)) + geom_point() + geom_line() +  
  scale_x_date(NULL, date_breaks = "1 month", date_labels = "%b")
```



State's school terms: summer break in 2013 was Jun 26–Sep 9.


The three school terms

```
term <- function(date) {  
  cut(date, breaks = ymd(20130101, 20130605, 20130825, 20140101),  
      labels = c("spring", "summer", "fall"))  
}  
daily <- daily %>% mutate(term = term(date))  
  
daily %>% filter(wday == "Sat") %>% ggplot(aes(date, n, colour = term)) +  
  geom_point(alpha = 1/3) + geom_line() +  
  scale_x_date(NULL, date_breaks = "1 month", date_labels = "%b")
```



School terms and day of week

```
daily %>%  
  ggplot(aes(wday, n, colour = term)) + geom_boxplot()
```



An improved model

```
mod1 <- lm(n ~ wday, data = daily)
mod2 <- lm(n ~ wday * term, data = daily)
```


```
daily %>% gather_residuals(without_term = mod1, with_term = mod2) %>%
  ggplot(aes(date, resid, colour = model)) + geom_line(alpha = 0.75)
```



What's going on here?


```
grid <- daily %>% data_grid(wday, term) %>%  
  add_predictions(mod2, "n")
```

```
ggplot(daily, aes(wday, n)) + geom_boxplot() +  
  geom_point(data = grid, colour = "red") + facet_wrap(~ term)
```



```
mod3 <- MASS::rlm(n ~ wday * term, data = daily)

daily %>% add_residuals(mod3, "resid") %>% ggplot(aes(date, resid)) +
  geom_hline(yintercept = 0, size = 2, colour = "white") + geom_line()
```



Either bundled up into a function:

```
compute_vars <- function(data) {  
  data %>%  
    mutate(term = term(date),  
           wday = wday(date, label = TRUE))  
}
```

Or directly in the model formula:

```
wday2 <- function(x) wday(x, label = TRUE)  
mod3 <- lm(n ~ wday2(date) * term(date), data = daily)
```


Time of year: an alternative approach COLUMBIA UNIVERSITY IN THE CITY OF NEW YORK

```
library(splines)
mod <- MASS::rlm(n ~ wday * ns(date, 5), data = daily)

daily %>% data_grid(wday, date = seq_range(date, n = 13)) %>%
  add_predictions(mod) %>%
  ggplot(aes(date, pred, colour = wday)) + geom_line() + geom_point()
```

