


# Lecture 8: Modeling II

Data Science for Business Analytics





Program

- This morning:
  - how models work mechanistically (focus on linear models),
  - how to use models to find patterns in real data.
- This afternoon:
  - ▶ how to use **many** simple models,
  - how to combine modeling and programming tools.

As usual, most of the material is borrowed from R for data science.

## Many models



### To work with large numbers of models, use

- 1. many simple models to better understand complex datasets.
- 2. list-columns to store arbitrary data structures in a data frame.
- 3. the **broom** package to turn models into tidy data.

### Note that this part:

- is harder than the others,
- and requires a deeper internalization of ideas (e.g., about modelling, data structures, and iteration).

### **Outline**



1 Iterations

2 gapminder

3 List-columns

## Reducing code duplication



#### Three main benefits:

- 1. Easier to see the intent of your code.
- 2. Easier to respond to changes in requirements.
- 3. Likely to have fewer bugs.

### Main tools for reducing duplication:

- Functions identify repeated patterns of code and extract them out into independent pieces.
- Iteration helps you when you need to do the same thing to multiple inputs.

## For loops



```
df <- tibble(a = rnorm(10), b = rnorm(10), c = rnorm(10), d = rnorm(10))</pre>
c(median(df$a), median(df$b), median(df$c), median(df$d))
## [1] 0.20443713 -0.35536298 0.09003127 -0.25752642
output <- vector("double", ncol(df)) # 1. output</pre>
for (i in seq_along(df)) {
                             # 2. sequence
  output[[i]] <- median(df[[i]]) # 3. body</pre>
output
## [1] 0.20443713 -0.35536298 0.09003127 -0.25752642
A side note:
v <- vector("double", 0)</pre>
seq_along(y)
## integer(0)
1:length(y)
## [1] 1 0
```

## Modifying an existing object



```
df <- tibble(a = rnorm(10),
             b = rnorm(10).
              c = rnorm(10),
             d = rnorm(10)
rescale01 <- function(x) {
    rng <- range(x, na.rm = TRUE)</pre>
    (x - rng[1]) / (rng[2] - rng[1])
}
df$a <- rescale01(df$a)
df$b <- rescale01(df$b)
df$c <- rescale01(df$c)
df$d <- rescale01(df$d)
for (i in seq_along(df)) {
    df[[i]] <- rescale01(df[[i]])</pre>
```

## For loops vs. functionals



### To compute the mean of every column:

```
output <- vector("double", length(df))
for (i in seq_along(df)) {
  output[[i]] <- mean(df[[i]])
}</pre>
```

#### As a function:

```
col_mean <- function(df) {
  output <- vector("double", length(df))
  for (i in seq_along(df)) {
    output[i] <- mean(df[[i]])
  }
  output
}</pre>
```

### How about other quantities?



```
col median <- function(df) {</pre>
  output <- vector("double", length(df))</pre>
  for (i in seq_along(df)) {
    output[i] <- median(df[[i]])</pre>
  output
col_sd <- function(df) {</pre>
  output <- vector("double", length(df))</pre>
  for (i in seq_along(df)) {
    output[i] <- sd(df[[i]])</pre>
  output
```

What's "wrong" here?

## What's "wrong" here?



```
f1 <- function(x) abs(x - mean(x)) ^ 1
f2 <- function(x) abs(x - mean(x)) ^ 2
f3 <- function(x) abs(x - mean(x)) ^ 3
```

## What's "wrong" here?



```
f1 <- function(x) abs(x - mean(x)) ^ 1
f2 <- function(x) abs(x - mean(x)) ^ 2
f3 <- function(x) abs(x - mean(x)) ^ 3
```

### Without duplication:

```
f <- function(x, i) abs(x - mean(x)) ^ i
```

### Back to column summaries



```
col_summary <- function(df, fun) {
  out <- vector("double", length(df))
  for (i in seq_along(df)) {
    out[i] <- fun(df[[i]])
  }
  out
}

col_summary(df, median)

## [1] 0.3684051 0.4523108 0.8108251 0.5468810

col_summary(df, mean)

## [1] 0.4406485 0.4485883 0.6825104 0.5249303</pre>
```

## **Functional programming**



- Use functions that return functions as output.
- Pass functions as arguments to others function.

### The purrr package:

- Functions eliminating the need for many common for loops.
- Similar to the apply family in base R (apply(), lapply(), tapply(), etc), but more consistent and easier to learn.

The goal is to break code into independent pieces:

- 1. Solve a problem for a single element of the list.
  - ▶ Once this is done, purrr generalizes to every element in the list.
- 2. Break a complex problem down into bite-sized pieces.
  - ▶ With purrr, small pieces are composed together with the pipe.

## The map functions



- map() makes a list.
- map\_lgl() makes a logical vector.
- map\_int() makes an integer vector.
- map\_dbl() makes a double vector.
- map\_chr() makes a character vector.

#### Each function:

- 1. Takes a vector as input.
- 2. Applies a function to each piece.
- 3. Returns a new vector that's the same length (and has the same names) as the input.

The return type is determined by the suffix.

## **Using map functions**



```
map_dbl(df, mean)
map_dbl(df, median)
map_dbl(df, sd)
##
## 0.4406485 0.4485883 0.6825104 0.5249303
##
## 0.3684051 0.4523108 0.8108251 0.5468810
##
                    b
          a
## 0.3362143 0.2640144 0.3516401 0.2929669
Using the pipe:
df %>% map_dbl(mean)
df %>% map_dbl(median)
df %>% map_dbl(sd)
##
## 0.4406485 0.4485883 0.6825104 0.5249303
##
          a
                    b
## 0.3684051 0.4523108 0.8108251 0.5468810
##
                    b
## 0.3362143 0.2640144 0.3516401 0.2929669
```



- All purrr functions are implemented in C (i.e., slightly faster).
- The second argument, .f, the function to apply, can be a formula, a character vector, or an integer vector.
- map\_\*() uses ... ([dot dot dot]) to pass along additional arguments to .f each time it's called:

■ The map functions preserve names:

```
z <- list(x = 1:3, y = 4:5)
map_int(z, length)
## x y
## 3 2</pre>
```



Splits the mtcars dataset into three pieces and fits the same linear model to each piece:

```
models <- mtcars %>%
  split(.$cyl) %>%
  map(function(df) lm(mpg ~ wt, data = df))
```

### Using a one-sided formula:

```
models <- mtcars %>%
split(.$cyl) %>%
map(~lm(mpg ~ wt, data = .))
```



### Extract a summary statistic like the $R^2$ :

### **Shortcut 3**



### Use an integer:

```
x <- list(list(1, 2, 3), list(4, 5, 6), list(7, 8, 9))
x %>% map_dbl(2)
```

## [1] 2 5 8

## Dealing with failure using safely()



- safely() is an adverb: it takes a function (a verb) and returns a modified version.
- The modified function always returns a list with two elements:
  - 1. result is the original result.
  - 2. error is an error object.

```
safe_log <- safely(log)</pre>
str(safe_log(10))
## List of 2
## $ result: num 2.3
## $ error : NULL
str(safe_log("a"))
## List of 2
    $ result: NULL
    $ error :List of 2
##
     ..$ message: chr "non-numeric argument to mathematical function"
     ..$ call : language log(x = x, base = base)
##
##
     ..- attr(*, "class")= chr [1:3] "simpleError" "error" "condition"
```




```
x \leftarrow list(1, 10, "a")
y <- x %>% map(safely(log))
str(y)
## List of 3
## $ :List of 2
## ..$ result: num 0
## ..$ error : NULL
## $ :List of 2
## ..$ result: num 2.3
## ..$ error : NULL
   $ :List of 2
##
## ..$ result: NULL
    ..$ error :List of 2
##
     .... $ message: chr "non-numeric argument to mathematical function"
##
     .... $ call : language log(x = x, base = base)
##
##
     ... - attr(*, "class") = chr [1:3] "simpleError" "error" "condition"
```

## transpose()



```
y <- y %>% transpose()
str(y)
## List of 2
    $ result:List of 3
##
    ..$ : num 0
   ..$ : num 2.3
##
##
    ..$ : NULL.
##
   $ error :List of 3
##
     ..$ : NULL
##
     ..$ : NULL.
##
     ..$ :List of 2
     .... $ message: chr "non-numeric argument to mathematical function"
##
     .... $ call : language log(x = x, base = base)
##
##
     ... - attr(*, "class")= chr [1:3] "simpleError" "error" "condition"
```





transpose(x)

## **Typical use**



```
is_ok <- y$error %>% map_lgl(is_null)
x[!is_ok]
## [[1]]
## [1] "a"
y$result[is_ok] %>% flatten_dbl()
## [1] 0.000000 2.302585
```

### Two other useful adverbs



possibly(): "simpler" than safely(), because you give it a default value to return when there is an error.

```
list(1, 10, "a") %>% map_dbl(possibly(log, NA_real_))
## [1] 0.000000 2.302585 NA
```

quietly(): instead of capturing errors, it captures printed output, messages, and warnings.

```
list(1, -1) %>% map(quietly(log)) %>% str()
## List of 2
## $ :List of 4
## ..$ result : num 0
    ..$ output : chr ""
    ..$ warnings: chr(0)
##
    ..$ messages: chr(0)
##
##
   $:List of 4
##
     ..$ result : num NaN
     ..$ output : chr ""
##
     .. $ warnings: chr "NaNs produced"
##
     ..$ messages: chr(0)
##
```

### Mapping over multiple arguments



```
mu \leftarrow list(5, 10, -3)
mu \%\% map(rnorm, n = 5) \%\% str()
## List of 3
## $ : num [1:5] 4.2 6.25 5.77 4.78 4.58
## $ : num [1:5] 9.58 11 9.72 11.26 10.65
## $ : num [1:5] -1.7 -3.87 -2.99 -3.88 -2.4
What if you also want to vary the standard deviation?
sigma <- list(1, 5, 10)
seq_along(mu) %>% map(~rnorm(5, mu[[.]], sigma[[.]])) %>% str()
## List of 3
## $ : num [1:5] 5.12 4.72 6.46 5.23 6
## $ : num [1:5] 13.91 6.12 6.92 10.23 4.35
## $ : num [1:5] 2.77 -15.81 13.25 -8.01 13.78
```



```
map2(mu, sigma, rnorm, n = 5) \%
    str()
## List of 3
   $ : num [1:5] 4.59 4.03 5.03 5.03 3.32
## $ : num [1:5] 15.3 4.4 11.7 12.5 10.7
##
   $ : num [1:5] -4.19 -1.02 -13.69 -11.03 -14.14
                     sigma
             mu
                              map2(mu, sigma, rnorm, n = 5)
                                   rnorm(5, 1, n = 5)
                                   rnorm(10, 5, n = 5)
             10
                       5
                                  rnorm(-3, 10, n = 5)
```

### Note that the arguments that

- vary for each call come before the function,
- are the same for every call come *after* (using ...).



```
n \leftarrow list(1, 3, 5)
args1 <- list(n, mu, sigma)</pre>
args1 %>%
    pmap(rnorm) %>%
    str()
## List of 3
## $ : num 6.58
## $ : num [1:3] 17.49 11.31 3.84
## $ : num [1:5] -3.04 12.12 -7.76 4.98 -12.74
                                              pmap(args1)
                      args1
                                            rnorm(1, 5, 1)
                        10
                                            rnorm(3, 10, 5)
                                           rnorm(5, -3, 10)
                                 10
```



```
args2 <- list(mean = mu, sd = sigma, n = n)</pre>
args2 %>%
    pmap(rnorm) %>%
    str()
## List of 3
    $: num 5.69
## $ : num [1:3] 5.22 3.84 5.22
##
    $ : num [1:5] -11.7 -12.11 4.41 -2.31 -6.24
                     args2
                                              pmap(args2)
                     sigma
               mu
                             n
                                      rnorm(mean = 5, sigma = 1, n = 1)
                             3
                                     rnorm(mean = 10, sigma = 5, n = 3)
               10
                                     rnorm(mean = -3, sigma = 10, n = 5)
                             5
```



```
params <- tribble(
 ~mean, ~sd, ~n,
   5, 1, 1,
  10, 5, 3,
  -3, 10, 5)
params %>%
 pmap(rnorm)
## [[1]]
## [1] 3.913497
##
## [[2]]
## [1] 4.920355 6.161049 4.401400
##
## [[3]]
## [1] -7.481742 1.717364 -14.804907 11.702570 -16.114206
```

## **Invoking different functions**



```
f <- c("runif", "rnorm", "rpois")</pre>
param \leftarrow list(list(min = -1, max = 1), list(sd = 5), list(lambda = 10))
invoke_map(f, param, n = 5) \%\% str()
## List of 3
    $ : num [1:5] -0.0769 -0.2496 0.9822 -0.6473 0.6269
## $ : num [1:5] -7.44 -5.38 5 -3.11 -6.92
## $ : int [1:5] 15 11 9 13 12
                             params
                                          invoke_map(f, params, n = 5)
                            min
                                max
                                         runif(min = -1, max = 1, n = 5)
                "runif"
                              sd
                "rnorm"
                                              rnorm(sd = 5, n = 5)
                             1ambda
               "rpois"
                                            rpois(lambda = 10, n = 5)
                              10
```

## Invoking different cont'd



```
sim <- tribble(~f, ~params,</pre>
               "runif", list(min = -1, max = 1),
               "rnorm", list(sd = 5).
               "rpois", list(lambda = 10))
sim %>%
   mutate(sim = invoke_map(f, params, n = 10)) %>%
    str()
## Classes 'tbl df', 'tbl' and 'data.frame': 3 obs. of 3 variables:
##
   $ f : chr "runif" "rnorm" "rpois"
##
   $ params:List of 3
## ..$ :List of 2
     ....$ min: num -1
##
##
     .. .. $ max: num 1
    ..$ :List of 1
##
##
     .. ..$ sd: num 5
##
    ..$ :List of 1
    .. .. $ lambda: num 10
##
##
    $ sim :List of 3
     ..$: num -0.464 0.524 0.973 -0.413 -0.201 ...
##
     ..$: num 1.038 11.54 0.529 2.285 -0.386 ...
##
##
     ..$: int 13 6 8 13 11 9 10 7 8 7
```

### **Outline**



1 Iterations

2 gapminder

3 List-columns

## gapminder

library(gapminder)



- Summarizes the progression of countries over time using variables like life expectancy and GDP.
- Popularized by Hans Rosling, a Swedish doctor and statistician, in a short video filmed in conjunction with the BBC


```
gapminder
## # A tibble: 1.704 x 6
##
      country
                  continent
                             year lifeExp
                                               pop gdpPercap
      <fct>
                  <fct>
                            <int>
                                    <dbl>
                                                       <db1>
##
                                             <int>
##
    1 Afghanistan Asia
                             1952
                                     28.8 8425333
                                                        779.
##
    2 Afghanistan Asia
                             1957
                                     30.3 9240934
                                                        821.
##
    3 Afghanistan Asia
                             1962
                                     32.0 10267083
                                                        853.
##
   4 Afghanistan Asia
                             1967
                                     34.0 11537966
                                                        836.
    5 Afghanistan Asia
                             1972
                                                        740.
##
                                     36.1 13079460
##
    6 Afghanistan Asia
                             1977
                                     38.4 14880372
                                                        786.
##
   7 Afghanistan Asia
                             1982
                                     39.9 12881816
                                                        978.
##
   8 Afghanistan Asia
                             1987
                                     40.8 13867957
                                                        852.
    9 Afghanistan Asia
##
                             1992
                                     41.7 16317921
                                                        649.
   10 Afghanistan Asia
                             1997
                                     41.8 22227415
                                                        635.
## # ... with 1.694 more rows
```

### Focus on three variables



How does life expectancy (lifeExp) change over time (year) for each country (country)?


```
gapminder %>%
ggplot(aes(year, lifeExp, group = country)) +
geom_line(alpha = 1/3)
```



## Model for a single country



```
nz <- filter(gapminder, country == "New Zealand")
nz %>% ggplot(aes(year, lifeExp)) + geom_line() + ggtitle("Full data = ")
nz_mod <- lm(lifeExp ~ year, data = nz)
nz %>% add_predictions(nz_mod) %>%
    ggplot(aes(year, pred)) + geom_line() + ggtitle("Linear trend + ")
nz %>% add_residuals(nz_mod) %>%
    ggplot(aes(year, resid)) +
    geom_hline(yintercept = 0, colour = "white", size = 3) +
    geom_line() + ggtitle("Remaining pattern")
```



### **Nested data**



```
(by_country <- gapminder %>%
  group_by(country, continent) %>%
  nest())
## # A tibble: 142 x 3
##
       country
                    continent data
       <fct>
                    <fct>
                                st.>
##
    1 Afghanistan Asia <tibble [12 x 4]>
##
    2 Albania
##
                    Europe <tibble [12 x 4]>
    3 Algeria Africa
                                <tibble \lceil 12 \times 4 \rceil >
##
##
    4 Angola
                Africa
                               \langle \text{tibble } [12 \times 4] \rangle
    5 Argentina Americas <tibble [12 x 4]>
##
    6 Australia Oceania
                                <tibble \lceil 12 \times 4 \rceil >
##
##
   7 Austria
                    Europe
                                \langle \text{tibble } [12 \times 4] \rangle
## 8 Bahrain
                    Asia
                                <tibble \lceil 12 \times 4 \rceil >
##
    9 Bangladesh Asia
                                <tibble \lceil 12 \times 4 \rceil >
## 10 Belgium
                    Europe
                                \langle \text{tibble } [12 \times 4] \rangle
## # ... with 132 more rows
```

- In a grouped data frame, each row is an observation.
- In a nested data frame, each row is a group.

#### **List-columns**



#### A model-fitting function applied to every country:

```
country_model <- function(df) lm(lifeExp ~ year, data = df)
models <- map(by_country$data, country_model)</pre>
```

#### Or add an additional list-column:

```
(by_country <- by_country %>% mutate(model = map(data, country_model)))
## # A tibble: 142 x 4
##
      country
                   continent data
                                                 model
##
      <fct> <fct> <fct> <fct> <fct> <fct> 
                                                 st>
##
    1 Afghanistan Asia <tibble [12 x 4]> <S3: lm>
    2 Albania
                   Europe \langle \text{tibble [12 x 4]} \rangle \langle \text{S3: lm} \rangle
##
##
    3 Algeria Africa <tibble [12 x 4]> <S3: lm>
    4 Angola Africa
                             <tibble [12 x 4]> <S3: lm>
##
    5 Argentina Americas
                              <tibble [12 x 4]> <S3: lm>
##
    6 Australia Oceania
##
                              <tibble [12 x 4] > <S3: lm>
   7 Austria Europe
                              <tibble [12 x 4]> <S3: lm>
##
## 8 Bahrain
                  Asia
                              <tibble [12 x 4] > <S3: lm>
    9 Bangladesh Asia
##
                              \langle \text{tibble } \lceil 12 \times 4 \rceil \rangle \langle \text{S3: } 1 \text{m} \rangle
## 10 Belgium
                              <tibble [12 x 4]> <S3: lm>
                   Europe
## # ... with 132 more rows
```

## Why bother?



- Avoid leaving the list of models as a free-floating object.
- No need to manually keep them in sync when using e.g. filter or arrange.

```
by_country %>% filter(continent == "Europe")
## # A tibble: 30 x 4
##
                             continent data
                                                         model
      country
##
      <fct>
                             <fct>
                                       st>
                                                         st>
    1 Albania
                                       <tibble [12 x 4]> <S3: lm>
##
                             Europe
##
    2 Austria
                             Europe
                                       <tibble [12 x 4]> <S3: lm>
                             Europe
                                       <tibble [12 x 4]> <S3: lm>
##
   3 Belgium
   4 Bosnia and Herzegovina Europe
                                       <tibble [12 x 4]> <S3: lm>
##
##
   5 Bulgaria
                             Europe
                                       <tibble [12 x 4]> <S3: lm>
   6 Croatia
                                       <tibble [12 x 4]> <S3: lm>
##
                             Europe
##
    7 Czech Republic
                             Europe
                                       <tibble [12 x 4] > <S3: lm>
##
   8 Denmark
                             Europe
                                       <tibble [12 x 4]> <S3: lm>
##
    9 Finland
                             Europe
                                       <tibble [12 x 4]> <S3: lm>
## 10 France
                                       <tibble [12 x 4]> <S3: lm>
                             Europe
## # ... with 20 more rows
```

## **Adding residuals**



```
(by_country <- by_country %>%
 mutate(resids = map2(data, model, add residuals)))
## # A tibble: 142 x 5
##
     country
                 continent data
                                            model
                                                     resids
     <fct>
##
                 <fct>
                           st>
                                            st>
                                                     st>
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
##
   1 Afghanistan Asia
   2 Albania
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
##
                 Europe
##
   3 Algeria
                 Africa
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
##
   4 Angola
               Africa
##
   5 Argentina Americas
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
   6 Australia
##
                 Oceania
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
   7 Austria
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
##
                 Europe
   8 Bahrain
                 Asia
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
##
##
   9 Bangladesh Asia
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
## 10 Belgium
                           <tibble [12 x 4]> <S3: lm> <tibble [12 x 5]>
                 Europe
## # ... with 132 more rows
```

## **Unnesting**

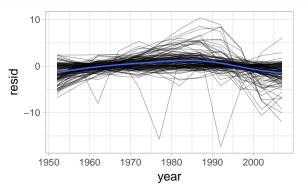


```
resids <- unnest(by_country, resids)
resids
## # A tibble: 1,704 x 7
##
                 continent
                                             pop gdpPercap
     country
                            year lifeExp
                                                            resid
##
     <fct>
                 <fct>
                           <int>
                                  <dbl>
                                           <int>
                                                     <dbl>
                                                             <dbl>
##
   1 Afghanistan Asia
                            1952
                                   28.8 8425333
                                                      779. -1.11
   2 Afghanistan Asia
                            1957
                                   30.3 9240934
##
                                                      821. -0.952
   3 Afghanistan Asia
                            1962
                                   32.0 10267083
                                                      853. -0.664
##
   4 Afghanistan Asia
##
                            1967
                                   34.0 11537966
                                                      836. -0.0172
   5 Afghanistan Asia
                            1972
                                   36.1 13079460
                                                      740.
                                                            0.674
##
##
   6 Afghanistan Asia
                            1977
                                   38.4 14880372
                                                      786. 1.65
   7 Afghanistan Asia
##
                            1982
                                   39.9 12881816
                                                      978. 1.69
##
   8 Afghanistan Asia
                            1987
                                   40.8 13867957
                                                      852. 1.28
   9 Afghanistan Asia
##
                            1992
                                   41.7 16317921
                                                      649. 0.754
## 10 Afghanistan Asia
                            1997
                                   41.8 22227415
                                                      635. -0.534
## # ... with 1.694 more rows
```

## Visualizing the residuals



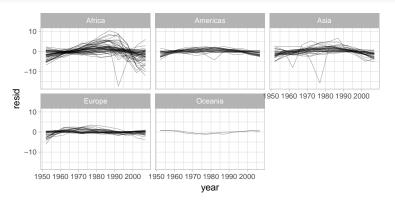
```
resids %>%
  ggplot(aes(year, resid)) +
  geom_line(aes(group = country), alpha = 1 / 3) +
  geom_smooth(se = FALSE)
```



# Visualizing the residuals cont'd



```
resids %>% ggplot(aes(year, resid, group = country)) +
   geom_line(alpha = 1 / 3) + facet_wrap(~continent)
```



## **Model quality**

## #



```
broom::glance(nz_mod)
##
    r.squared adj.r.squared sigma statistic p.value df
## 1 0.9535846
                0.9489431 0.8043472 205.4459 5.407324e-08 2
##
      logLik
                 AIC
                      BIC deviance df.residual
## 1 -13.32064 32.64128 34.096 6.469743
                                          10
by_country %>% mutate(glance = map(model, glance)) %>% unnest(glance)
## # A tibble: 142 x 16
##
     country continent data model resids r.squared adj.r.squared sigma
     <fct> <fct>
##
                    <dbl>
                                                    <dbl> <dbl>
   1 Afghan~ Asia <tib~ <S3:~ <tibb~
                                        0.948
                                                    0.942 1.22
##
##
   2 Albania Europe <tib~ <S3:~ <tibb~ 0.911
                                                    0.902 1.98
   3 Algeria Africa <tib~ <S3:~ <tibb~ 0.985
                                                    0.984 1.32
##
   4 Angola Africa <tib~ <S3:~ <tibb~ 0.888
##
                                                    0.877 1.41
##
   5 Argent~ Americas <tib~ <S3:~ <tibb~
                                        0.996
                                                    0.995 0.292
##
   6 Austra~ Oceania
                   <tib~ <S3:~ <tibb~ 0.980
                                                    0.978 0.621
##
  7 Austria Europe <tib~ <S3:~ <tibb~ 0.992
                                                    0.991 0.407
##
   8 Bahrain Asia <tib~ <S3:~ <tibb~ 0.967
                                                    0.963 1.64
   0.988 0.977
##
## 10 Belgium Europe <tib~ <S3:~ <tibb~ 0.995
                                                    0.994 0.293
## # ... with 132 more rows, and 8 more variables: statistic <dbl>...
## # p.value <dbl>, df <int>, logLik <dbl>, AIC <dbl>, BIC <dbl>,
```

deviance <dbl>, df.residual <int>

### Or better



```
glance <- by_country %>%
 mutate(glance = map(model, broom::glance)) %>%
 unnest(glance, .drop = TRUE)
glance
## # A tibble: 142 x 13
     country continent r.squared adj.r.squared sigma statistic p.value
##
##
     <fct>
             <fct>
                          <dbl>
                                        <dbl> <dbl>
                                                       <dbl>
                                                                <dbl>
##
   1 Afghan~ Asia
                          0.948
                                        0.942 1.22
                                                       181. 9.84e- 8
   2 Albania Europe
                                        0.902 1.98
                                                       102. 1.46e- 6
##
                          0.911
##
   3 Algeria Africa
                          0.985
                                        0.984 1.32
                                                       662.
                                                             1.81e-10
   4 Angola Africa
                          0.888
                                        0.877 1.41
                                                        79.1 4.59e- 6
##
##
   5 Argent~ Americas
                          0.996
                                        0.995 0.292
                                                      2246. 4.22e-13
##
   6 Austra~ Oceania
                          0.980
                                        0.978 0.621
                                                       481. 8.67e-10
                                                      1261. 7.44e-12
##
   7 Austria Europe
                          0.992
                                        0.991 0.407
##
   8 Bahrain Asia
                          0.967
                                        0.963 1.64
                                                       291. 1.02e- 8
##
   9 Bangla~ Asia
                          0.989
                                        0.988 0.977
                                                       930. 3.37e-11
## 10 Belgium Europe
                          0.995
                                        0.994 0.293
                                                      1822.
                                                             1.20e-12
## # ... with 132 more rows, and 6 more variables: df <int>,
## #
      logLik <dbl>, AIC <dbl>, BIC <dbl>, deviance <dbl>,
      df.residual <int>
## #
```

#### Which models don't fit well?



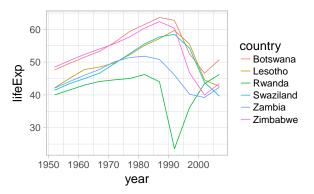
```
glance %>%
 arrange(r.squared)
## # A tibble: 142 x 13
              continent r.squared adj.r.squared sigma statistic p.value
##
     country
##
     <fct>
              <fct>
                           dbl>
                                        <dbl> <dbl>
                                                       <dbl>
                                                              <dbl>
                                     -0.0811 6.56
                                                       0.175 0.685
##
   1 Rwanda Africa
                          0.0172
   2 Botswana Africa
                          0.0340
                                     -0.0626 6.11
                                                       0.352
                                                             0.566
##
                                     -0.0381 7.21
##
   3 Zimbabwe Africa
                          0.0562
                                                       0.596 0.458
   4 Zambia
             Africa
                          0.0598
                                     -0.0342 4.53
                                                       0.636 0.444
##
##
   5 Swazila~ Africa
                          0.0682
                                     -0.0250 6.64
                                                       0.732
                                                             0.412
##
   6 Lesotho Africa
                          0.0849
                                     -0.00666 5.93
                                                       0.927
                                                             0.358
##
   7 Cote d'~ Africa
                          0.283
                                      0.212
                                              3.93
                                                       3.95
                                                             0.0748
##
   8 South A~ Africa
                          0.312
                                      0.244 4.74
                                                       4.54 0.0588
                                              3.19
                                                       5.20
                                                             0.0457
##
   9 Uganda
             Africa
                          0.342
                                      0.276
  10 Congo. ~ Africa
                          0.348
                                      0.283
                                              2.43
                                                       5.34
                                                             0.0434
## #
    ... with 132 more rows, and 6 more variables: df <int>,
## #
      logLik <dbl>, AIC <dbl>, BIC <dbl>, deviance <dbl>,
## #
      df.residual <int>
```





```
bad_fit <- filter(glance, r.squared < 0.25)

gapminder %>% semi_join(bad_fit, by = "country") %>%
    ggplot(aes(year, lifeExp, colour = country)) + geom_line()
```



### **Outline**



1 Iterations

2 gapminder

3 List-columns

#### **List-columns**



```
tibble(x = list(1:3, 3:5),
      y = c("1, 2", "3, 4, 5"))
## # A tibble: 2 x 2
## x
             У
## <list> <chr>
## 1 <int [3]> 1, 2
## 2 <int [3] > 3, 4, 5
tribble(~x, ~y,
      1:3, "1, 2",
       3:5, "3, 4, 5")
## # A tibble: 2 x 2
## x
## <list> <chr>
## 1 <int [3] > 1, 2
## 2 <int [3] > 3, 4, 5
```

# An effective list-column pipeline



- Create the list-column using one of nest(), summarise() + list(), or mutate() + a map function.
- 2. Create other intermediate list-columns by transforming existing list columns with map(), map2() or pmap().
- Simplify the list-column back down to a data frame or atomic vector.

## **Creating list-columns**



- 1. With nest() to convert a grouped data frame into a nested data frame where you have list-column of data frames.
- 2. With mutate() and vectorised functions that return a list.
- With summarise() and summary functions that return multiple results.

## Create with nesting I



```
gapminder %>% group_by(country, continent) %>% nest()
## # A tibble: 142 x 3
##
       country
                      continent data
##
       <fct>
                      <fct>
                                  st>
     1 Afghanistan Asia <tibble [12 x 4]>
##
                      Europe <tibble [12 x 4]>
##
     2 Albania
##
     3 Algeria
                      Africa
                                  <tibble \lceil 12 \times 4 \rceil >
##
     4 Angola
                      Africa
                                  <tibble \lceil 12 \times 4 \rceil >
##
     5 Argentina Americas
                                  \langle \text{tibble } [12 \times 4] \rangle
    6 Australia Oceania
                                  <tibble \lceil 12 \times 4 \rceil >
##
##
     7 Austria
                                  <tibble [12 x 4]>
                      Europe
                                  \langle \text{tibble } \lceil 12 \times 4 \rceil \rangle
##
    8 Bahrain
                      Asia
     9 Bangladesh Asia
                                  <tibble \lceil 12 \times 4 \rceil >
##
## 10 Belgium
                      Europe
                                  \langle \text{tibble } [12 \times 4] \rangle
## # ... with 132 more rows
```

## Create with nesting II



```
gapminder %>% nest(year:gdpPercap)
## # A tibble: 142 x 3
##
                      continent data
       country
##
       <fct>
                      <fct>
                                   st>
     1 Afghanistan Asia
                                  <tibble [12 x 4]>
##
                      Europe <tibble [12 x 4]>
##
     2 Albania
##
     3 Algeria
                      Africa
                                   <tibble \lceil 12 \times 4 \rceil >
##
     4 Angola
                      Africa
                                   <tibble \lceil 12 \times 4 \rceil >
##
     5 Argentina Americas
                                   \langle \text{tibble } [12 \times 4] \rangle
    6 Australia
                      Oceania
                                   <tibble \lceil 12 \times 4 \rceil >
##
##
     7 Austria
                                   <tibble [12 x 4]>
                      Europe
                                   \langle \text{tibble } \lceil 12 \times 4 \rceil \rangle
##
    8 Bahrain
                      Asia
     9 Bangladesh Asia
                                   <tibble \lceil 12 \times 4 \rceil >
##
## 10 Belgium
                      Europe
                                   \langle \text{tibble } [12 \times 4] \rangle
## # ... with 132 more rows
```

### **Create from vectorized functions**



```
df <- tribble(~x1, "a,b,c", "d,e,f,g")</pre>
df %>% mutate(x2 = stringr::str_split(x1, ","))
## # A tibble: 2 x 2
## x1 x2
## <chr> <list>
## 1 a,b,c <chr [3]>
## 2 d,e,f,g <chr [4]>
sim <- tribble(~f, ~params,</pre>
              "runif", list(min = -1, max = -1),
              "rnorm", list(sd = 5),
              "rpois", list(lambda = 10))
sim %>% mutate(sims = invoke_map(f, params, n = 10))
## # A tibble: 3 x 3
## f params sims
## <chr> <list> <list>
## 1 runif <list [2]> <dbl [10]>
## 2 rnorm <list [1]> <dbl [10]>
## 3 rpois <list [1]> <int [10]>
```



#### What's wrong here?

## 2 6. <dbl [5]> ## 3 8. <dbl [5]>

```
mtcars %>% group_by(cyl) %>% summarise(q = quantile(mpg))
## Error in summarise_impl(.data, dots): Column `q` must be length 1 (a summary
Use list-columns:
mtcars %>% group_by(cyl) %>% summarise(q = list(quantile(mpg)))
## # A tibble: 3 x 2
##
       cyl q
     <dbl> t>
##
## 1 4. <dbl [5]>
```

### Create from multivalued summaries | | \( \psi \) Columbia University



```
probs \leftarrow c(0.01, 0.25, 0.5, 0.75, 0.99)
mtcars %>% group_by(cyl) %>%
  summarise(p = list(probs), q = list(quantile(mpg, probs))) %>%
  unnest()
    A tibble: 15 x 3
##
        cyl
              р
```

# **Simplifying list-columns**



- If you want a single value, use mutate() with map\_lgl(), map\_int(), map\_dbl(), and map\_chr() to create an atomic vector.
- 2. If you want many values, use unnest() to convert list-columns back to regular columns, repeating the rows as many times as necessary.

### List to vector

## 3 <db1 [5]> double





When columns contain different number of elements:

## Error: All nested columns must have the same number of elements.

# Making tidy data with broom



- 1. broom::glance(model)
  - A row for each model.
  - Columns give a model summary (measure of model quality, complexity, or combination of both).
- 2. broom::tidy(model)
  - A row for each coefficient in the model.
  - Columns give information about the estimate or its variability.
- 3. broom::augment(model, data)
  - A row for each row in data.
  - Adds extra values like residuals, and influence statistics.